

Red Hat Openshift 29 June to 3rd July 2026

Course designed and delivered by Jeganathan Swaminathan

Jeganathan Swaminathan

`jegan@tektutor.org`

`https://www.tektutor.org`

GitHub: `https://github.com/tektutor`

LinkedIn: `www.linkedin.com/in/jeganathan-swaminathan-2a6a6a6`

Day 1


Info - Hypervisor Overview

- nothing but virtualization
- with virtualization/hypervisor software, we can run multiple OS side by side on the same laptop/desktop/workstation/server
- there are 2 types of Hypervisors
 1. Type 1
 - a.k.a Baremetal Hypervisor
 - is used in Workstation & Server
 - performance wise it offers almost near native performance (3% less performance as opposed to running a OS on direct H/W)
 - this can be installed directly on top of a H/W with no Operating System
 - examples
 1. VMWare vSphere(vCenter) - Paid
 2. Linux KVM (opensource & Free)
 3. Microsoft Hyper-V
 2. Type 2
 - a.k.a Hosted Hypervisor
 - this can be installed only top of a Host Operating System (Windows, Mac OSX, Linux, etc)
 - is used in Workstation, Desktops & Laptops
 - examples
 - Oracle VirtualBox (Free - Windows/Linux/Mac)
 - VMWare Workstation (Free - after Broadcom acquired VMWare - Windows, Linux, Mac)
 - VMWare Fusion (Mac OS-X - Need a commercial license)
 - Parallels (Mac OS-X - Need a commercial license)
- virtualization helps organization save cost in many fold
 - with less physical servers, many virtual machines(Guest OS) can be supported
 - saves cost in terms of procuring less physical servers
 - saves cost in terms of power consumption
 - saves cost in terms of Air Condition
 - saves cost in terms of Sound proofing
 - saves cost in terms of Real estate rental/leasing
- this type of virtualization is called heavy-weight
 - because each VM requires dedicated Hardware resources
 - each VM runs a fully functional OS with its own dedicated OS Kernel
- each VM, represents one fully function Operating System


Info - Containerization

- application virtualization technology
- it is light weight, because each container represents a single application not an Operating System
- all the containers that runs on top a OS shares the Hardware resources available on the underlying OS
- containers will never be able to replace an Operating System or Virtualization
- in fact, in real-world scenarios, virtualization and containerization are used in combination
- it is a linux technology
- linux kernel features that enables containerization
 1. Namespace
 - helps isolation of containers from each other
 2. Control Groups (CGroups)
 - it helps us apply resource quota restrictions like
 - we can restrict how much RAM, disk and cpu can be utilized by a container
- container runtime
 - is a low-level software that manages container images and containers
 - it is not user-friendly, hence end-user like us never use container runtime directly
 - examples
 - runC
 - cRun
 - CRI-O
- container engine
 - is a high-level software that manages container images and containers
 - it is very user-friendly, but they depend on container runtimes internally to manage images and containers
 - examples
 - docker
 - podman

Info - Hypervisor High-Level Architecture



Info - Docker High-Level Architecture




Info - Docker image

- is a blueprint of a docker container
- you can imagine docker image similar to Window12.iso (OS Image) or Ubuntu24.04.iso(OS Image)
- with the help of a docker image, we can create any number of docker

containers

- docker images => typically has one application and its dependencies (libraries, dependent tools, etc.,)
- docker images are conservatively built, which means it only contains bare minimum tools required to run a specific application
- every docker images gets a unique name and unique id (256-bit HASH)



Info - Docker Container

- is the running instance of a Docker Image (Container Image)
- each container gets a unique name and ID (SHA Hash)
- each container gets its own dedicated software defined network stack (7 OSI Layers)
- each container get its own file system (files & folders)
- each container uses about 7 namespaces
- each container gets one or more IP addresses (generally Private IP)

Info - Docker Registry

- is a collection of multiple docker images
- there are 3 types

1. Local Docker Registry
 - is a folder created and maintained by Docker server on the same machine where it runs
 - for example, in linux distros, local docker registry is maintained under directory `/var/lib/docker`
2. Remote docker registry
 - is a web site (`hub.docker.com`) maintained by Docker Inc and opensource community
3. Private Docker registry
 - can be setup using JFrog Artifactory or Sonatype Nexus

Demo - Installing Docker community edition (Docker CE)


```
# Add Docker's official GPG key:
sudo apt update
sudo apt install ca-certificates curl
sudo install -m 0755 -d /etc/apt/keyrings
sudo curl -fsSL https://download.docker.com/linux/ubuntu/gpg -o /etc/apt/
keyrings/docker.asc
sudo chmod a+r /etc/apt/keyrings/docker.asc

# Add the repository to Apt sources:
sudo tee /etc/apt/sources.list.d/docker.sources <<EOF
Types: deb
URIs: https://download.docker.com/linux/ubuntu
Suites: $(. /etc/os-release && echo "${UBUNTU_CODENAME:-$VERSION_CODENAME}")
Components: stable
Architectures: $(dpkg --print-architecture)
Signed-By: /etc/apt/keyrings/docker.asc
EOF

sudo apt update
sudo apt install docker-ce docker-ce-cli containerd.io docker-buildx-plugin
docker-compose-plugin -y
sudo systemctl enable docker
sudo systemctl start docker
sudo usermod -aG docker $USER
newgrp docker
```


Lab - Checking docker details

```
docker --version
docker info
```

A terminal window titled 'jegan@tektutor: ~' showing the output of 'docker --version' and 'docker info'. The 'docker --version' command returns 'Docker version 29.2.1, build a5c7197'. The 'docker info' command returns detailed information about the Docker client, plugins, and server. The server information includes 3 containers (0 running, 0 paused, 3 stopped), 31 images, and various configuration details like storage driver (overlay2) and logging driver (json-file).

```
jegan@tektutor.org:~$ docker --version
Docker version 29.2.1, build a5c7197
jegan@tektutor.org:~$ docker info
Client: Docker Engine - Community
 Version: 29.2.1
 Context: default
 Debug Mode: false
 Plugins:
  buildx: Docker Buildx (Docker Inc.)
    Version: v0.31.1
    Path: /usr/libexec/docker/cli-plugins/docker-buildx
  compose: Docker Compose (Docker Inc.)
    Version: v5.0.2
    Path: /usr/libexec/docker/cli-plugins/docker-compose

Server:
 Containers: 3
  Running: 0
  Paused: 0
  Stopped: 3
 Images: 31
 Server Version: 29.2.1
 Storage Driver: overlay2
  Backing Filesystem: extfs
  Supports d_type: true
  Using metacopy: false
  Native Overlay Diff: true
 userxattr: false
 Logging Driver: json-file
 Cgroup Driver: systemd
```

A terminal window titled 'jegan@tektutor: ~' showing the continuation of the 'docker info' output. It lists runtime information (io.containerd.runc.v2, runc), security options (apparmor, seccomp), kernel version (7.0.0-22-generic), operating system (Ubuntu 26.04 LTS), architecture (x86_64), and hardware details (24 CPUs, 61.23GiB memory). It also shows the Docker root directory as '/var/lib/docker' and other configuration details like debug mode, username, and registries.

```
Runtimes: io.containerd.runc.v2 runc
Default Runtime: runc
Init Binary: docker-init
 containerd version: dea7da592f5d1d2b7755e3a161be07f43fad8f75
  runc version: v1.3.4-0-gd6d73eb8
   init version: de40ad0
Security Options:
 apparmor
 seccomp
   Profile: builtin
 cgroupns
Kernel Version: 7.0.0-22-generic
Operating System: Ubuntu 26.04 LTS
OSType: linux
Architecture: x86_64
CPUs: 24
Total Memory: 61.23GiB
Name: tektutor.org
ID: 9554324f-6d2b-40e7-8920-5e553fa0f8a3
Docker Root Dir: /var/lib/docker
Debug Mode: false
Username: tektutor
Experimental: false
Insecure Registries:
 ::1/128
 127.0.0.0/8
Live Restore Enabled: false
Firewall Backend: iptables

jegan@tektutor.org:~$
```

Lab - Listing docker images from your local docker registry

```
docker images
```

```
jegan@tektutor.org:~$ docker images
```

IMAGE	ID	DISK USAGE	CONTENT SIZE	EXTRA
alpine:latest	3cb067eab609	8.45MB	0B	
badimage-ubuntu:1.0	5736020bd49c	992MB	0B	
buchgr/bazel-remote-cache:latest	aad72c4b8fe1	38.5MB	0B	
database:latest	27deb4ec2bb4	445MB	0B	
docker/docker-bench-security:latest	0037349aef7e	51.6MB	0B	
gaiaadm/pumba:latest	fc4eb3c7aae3	28.9MB	0B	
gcr.io/cadvisor/cadvisor:v0.49.1	c02cf39d3dba	80.8MB	0B	
gcr.io/k8s-minikube/kicbase:v0.0.50	6da180ef5035	1.37GB	0B	U
goodimage-ubuntu:1.0	5ec1158c04f9	339MB	0B	
high-throughput-app:latest	da954fb959a3	62.3MB	0B	
high-throughput-app:v1.0	da954fb959a3	62.3MB	0B	
high-throughput-app:v2.0	da954fb959a3	62.3MB	0B	
high-throughput-app:v2.1	da954fb959a3	62.3MB	0B	
legacy-app:latest	da954fb959a3	62.3MB	0B	
legacy-app:v1.0	da954fb959a3	62.3MB	0B	
legacy-app:v2.0	da954fb959a3	62.3MB	0B	
legacy-app:v2.1	da954fb959a3	62.3MB	0B	
monitoring-agent:latest	da954fb959a3	62.3MB	0B	
monitoring-agent:v1.0	da954fb959a3	62.3MB	0B	
monitoring-agent:v2.0	da954fb959a3	62.3MB	0B	
monitoring-agent:v2.1	da954fb959a3	62.3MB	0B	
myapi:latest	da954fb959a3	62.3MB	0B	
myapi:v1.0	da954fb959a3	62.3MB	0B	
myapi:v2.0	da954fb959a3	62.3MB	0B	
myapi:v2.1	da954fb959a3	62.3MB	0B	
myapp:latest	da954fb959a3	62.3MB	0B	
myapp:modified	4ef0d95a33d4	62.3MB	0B	
myapp:v1.0	da954fb959a3	62.3MB	0B	

Lab - Downloading docker images from docker hub to local registry

```
docker pull hello-world:latest
docker pull ubuntu:latest

docker images
```

```
Mon Jun 29 2:09 PM
jegan@tektutor:~$ # Download docker images from docker hub remote registry to local docker registry
jegan@tektutor.org:~$ docker pull hello-world:latest
latest: Pulling from library/hello-world
4f55086f7dd0: Pull complete
Digest: sha256:96498ffd522e70807ab6384a5c0485a79b9c7c08ca79ba08623edcad1054e62d
Status: Downloaded newer image for hello-world:latest
docker.io/library/hello-world:latest
jegan@tektutor.org:~$
jegan@tektutor.org:~$ docker pull ubuntu:latest
latest: Pulling from library/ubuntu
81e2f2053c8f: Pull complete
d1f56e4c7f2f: Pull complete
Digest: sha256:53958ec7b67c2c9355df922dd08dbf0360611f8c3cdb656875e81873db9ffdba
Status: Downloaded newer image for ubuntu:latest
docker.io/library/ubuntu:latest
jegan@tektutor.org:~$
jegan@tektutor.org:~$ docker images | grep ubuntu
WARNING: This output is designed for human readability. For machine-readable output, please use --format.
badimage-ubuntu:1.0          5736020bd49c          992MB          0B
goodimage-ubuntu:1.0         5ec1158c04f9          339MB          0B
tektutor/ubuntu:1.0          0ffe51b3f464          1.08GB         0B
ubuntu:22.04                 86f1a8d7b38e          78.1MB         0B
ubuntu:latest                5a4c6b929c57          100MB          0B
jegan@tektutor.org:~$
```

Lab - Deleting a docker image from your local registry

```
docker images | grep hello
docker rmi hello-world:latest
docker images | grep hello
```

```
Mon Jun 29 2:21 PM
jegan@tektutor:~$ docker images | grep hello
WARNING: This output is designed for human readability. For machine-readable output, please use --format.
hello-world:latest          e2ac70e7319a          10.1kB         0B
tektutor/hello-microservice:1.0 540ee094e2c8          412MB          0B
tektutor/hello-ms:1.0         c588ed117d2c          411MB          0B
jegan@tektutor.org:~$
jegan@tektutor.org:~$ # Deleting a docker image from your local registry
jegan@tektutor.org:~$ docker rmi hello-world:latest
Untagged: hello-world:latest
Untagged: hello-world@sha256:96498ffd522e70807ab6384a5c0485a79b9c7c08ca79ba08623edcad1054e62d
Deleted: sha256:e2ac70e7319a02c5a477f5825259bd118b94e8b02c279c67afa63adab6d8685b
Deleted: sha256:897b3f2a7c1bc2f3d02432f7892fe31c6272c521ad4d70257df624504a3238b4
jegan@tektutor.org:~$ docker images | grep hello
WARNING: This output is designed for human readability. For machine-readable output, please use --format.
tektutor/hello-microservice:1.0 540ee094e2c8          412MB          0B
tektutor/hello-ms:1.0         c588ed117d2c          411MB          0B
jegan@tektutor.org:~$
```

Lab - Creating a container in the interactive(foreground) mode

From terminal 1


```
docker run -it --name ubuntu1-jegan --hostname ubuntu1-jegan ubuntu:latest /bin/
bash
hostname
hostname -i
ls
exit
```

List all running containers from second terminal

```
docker ps | grep jegan
```

List all running containers you created

```
docker ps -a | grep jegan
```



```
jegan@tektutor.org:~/custom-docker-image$ docker images | grep ubuntu:26.04
WARNING: This output is designed for human readability. For machine-readable output, please use --format.
ubuntu:26.04          5a4c6b929c57      100MB      0B
jegan@tektutor.org:~/custom-docker-image$
jegan@tektutor.org:~/custom-docker-image$ docker run -it --name ubuntu1 --hostname ubuntu1 ubuntu:latest /bin/bash
root@ubuntu1:/#
root@ubuntu1:/# hostname
ubuntu1
root@ubuntu1:/# hostname -i
172.17.0.2
root@ubuntu1:/# ls
bin boot dev etc home lib lib64 media mnt opt proc root run sbin srv sys tmp usr var
root@ubuntu1:/# exit
exit
jegan@tektutor.org:~/custom-docker-image$ docker ps
CONTAINER ID   IMAGE      COMMAND                  CREATED        STATUS        PORTS        NAMES
jegan@tektutor.org:~/custom-docker-image$ docker ps -a
CONTAINER ID   IMAGE      COMMAND                  CREATED        STATUS        PORTS        NAMES
8118792f4e8e   ubuntu:latest   "/bin/bash"            19 seconds ago   Exited (0) 7 seconds ago
a064e9819ff3   gcr.io/k8s-minikube/kicbase:v0.0.50   "/usr/local/bin/entr..." 27 hours ago     Exited (130) 23 hours ago
7cdbf105ddc8   30ba44506a6d   "/bin/bash"            10 days ago     Exited (137) 9 days ago
jegan@tektutor.org:~/custom-docker-image$
```

Lab - Create a custom docker image

In your home directory, create a folder something like custom-docker-image

```
cd ~/
mkdir -p custom-docker-image
cd custom-docker-image
touch Dockerfile
```

Add the below content to your Dockerfile, you can run gedit to open windows notepad like editor

```
FROM ubuntu:26.04

RUN apt update && apt install -y net-tools iputils-ping
RUN apt update && apt install -y default-jdk maven
```

Build your custom docker image

```
cd ~/custom-docker-image
docker build -t ubuntu-jegan:1.0 .
docker images | grep jegan
```

Create a container using your custom docker image

```
docker run -dit --name c1-jegan --hostname c1-jegan ubuntu-jegan:1.0 /bin/bash
docker exec -it c1-jegan /bin/bash
hostname
hostname -i
ping 8.8.8.8
ifconfig
ls
exit
```

List and see the container should still run because we created this container in background

```
docker ps
```

Lab - Stopping a running container

```
docker ps | grep c1-jegan
docker stop c1-jegan
```

```
docker ps | grep c1-jegan
docker ps -a | grep c1-jegan
```

Lab - Start an exited container

```
docker ps -a | grep c1-jegan
docker start c1-jegan
docker ps | grep c1-jegan
```

Lab - Restart a running container

```
docker ps a | grep c1-jegan
docker restart c1-jegan
docker ps | grep c1-jegan
```

Lab - Deleting a running container

In order to delete a running container gracefully, you must stop it first

```
docker stop c1-jegan
docker rm c1-jegan
docker ps -a
```

In order to delete a running container forcibly

```
docker rm -f c1-jegan
docker ps -a
```

Lab - Rename a container

```
docker run -dit ubuntu:latest /bin/bash
docker ps a
docker rename zealous_montalcini c2-jegan
docker ps
```



```
jegan@tektutor.org:~$ docker ps
CONTAINER ID   IMAGE     COMMAND   CREATED   STATUS    PORTS   NAMES
jegan@tektutor.org:~$ docker run -dit --name c1-jegan --hostname c1-jegan ubuntu:latest /bin/bash
5a8c2b4e46924008351d1d41c4bcd51544b0b262c69ae46573729ad094ca2f9f
jegan@tektutor.org:~$
jegan@tektutor.org:~$ docker run -dit ubuntu:latest /bin/bash
f93b087d7f6d8f27d5dc2248832f4809e82375fe7702ff389d1e6a3e8ccdf496
jegan@tektutor.org:~$ docker ps
CONTAINER ID   IMAGE     COMMAND   CREATED   STATUS    PORTS   NAMES
f93b087d7f6d   ubuntu:latest   "/bin/bash"   5 seconds ago   Up 4 seconds   zealous_montalcini
5a8c2b4e4692   ubuntu:latest   "/bin/bash"   2 minutes ago   Up 2 minutes   c1-jegan
jegan@tektutor.org:~$ docker rename zealous_montalcini c2-jegan
jegan@tektutor.org:~$ docker ps
CONTAINER ID   IMAGE     COMMAND   CREATED   STATUS    PORTS   NAMES
f93b087d7f6d   ubuntu:latest   "/bin/bash"   26 seconds ago   Up 25 seconds   c2-jegan
5a8c2b4e4692   ubuntu:latest   "/bin/bash"   2 minutes ago   Up 2 minutes   c1-jegan
jegan@tektutor.org:~$ docker exec -it c2-jegan /bin/bash
root@f93b087d7f6d:/# exit
exit
jegan@tektutor.org:~$ hostname
tektutor.org
jegan@tektutor.org:~$ docker exec -it c2-jegan hostname
f93b087d7f6d
jegan@tektutor.org:~$ docker exec -it c2-jegan hostname -i
172.17.0.3
jegan@tektutor.org:~$
```

Lab - Finding details about a container image

```
docker image inspect ubuntu:26.04
```

```
jegan@tektutor.org:~$ # Finding details about a container image
jegan@tektutor.org:~$ docker image inspect ubuntu:26.04
[
  {
    "Id": "sha256:5a4c6b929c57abe310fad22db2820f1423e645cdc9344bb05adaca9b50c3403f",
    "RepoTags": [
      "ubuntu:26.04",
      "ubuntu:latest"
    ],
    "RepoDigests": [
      "ubuntu@sha256:53958ec7b67c2c9355df922dd08dbf0360611f8c3cdb656875e81873db9ffdba"
    ],
    "Comment": "Add rock control metadata",
    "Created": "2026-06-10T03:29:33.588447386Z",
    "Config": {
      "Env": [
        "PATH=/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin"
      ],
      "Cmd": [
        "/bin/bash"
      ],
      "Labels": {
        "org.opencontainers.image.created": "2026-06-10T03:30:57.931695+00:00",
        "org.opencontainers.image.description": "The Ubuntu container image maintained by Canonical\n\nUbuntu is a Debian-based Linux operating system that runs from the desktop to the cloud, to all your internet connected things.\n\nIt is the world's most popular operating system across public clouds and OpenStack clouds.\n\nIt is the number one platform for containers; from Docker to Kubernetes to LXD, Ubuntu can run your containers at scale.\n\nFast, secure and simple, Ubuntu powers millions of PCs worldwide.\n",
        "org.opencontainers.image.title": "ubuntu",
        "org.opencontainers.image.version": "26.04"
      }
    }
  ]
}
```

```
    },
    "Architecture": "amd64",
    "Os": "linux",
    "Size": 100154952,
    "GraphDriver": {
      "Data": {
        "LowerDir": "/var/lib/docker/overlay2/6cf3ac66468bf0f3ea9d4349f47a5d70e90f6e10508e4a6ded47024c3fa4a4cb/diff",
        "MergedDir": "/var/lib/docker/overlay2/9a98a26589c570a1d29fc482219e1c1d455286d65d619a9003602668b7d97adb/merged",
        "UpperDir": "/var/lib/docker/overlay2/9a98a26589c570a1d29fc482219e1c1d455286d65d619a9003602668b7d97adb/diff",
        "WorkDir": "/var/lib/docker/overlay2/9a98a26589c570a1d29fc482219e1c1d455286d65d619a9003602668b7d97adb/work"
      }
    },
    "Name": "overlay2"
  },
  "RootFS": {
    "Type": "layers",
    "Layers": [
      "sha256:e8c084c1b320c172e8be941d735d77298e49986014213a8282c9b533ee216a61",
      "sha256:ab0ae19b58df847c6f0e4beb31f759a4dac2cc42288a6bda6f467eea4584c541"
    ]
  },
  "Metadata": {
    "LastTagTime": "0001-01-01T00:00:00Z"
  }
}
jegan@tektutor.org:~$
```

Lab - Finding details about a container

```
docker container inspect c1-jegan
docker inspect c1-jegan
```

```
Mon Jun 29 4:43 PM
jegan@tektutor:~
jegan@tektutor.org:~$ docker ps
CONTAINER ID   IMAGE          COMMAND                  CREATED        STATUS        PORTS        NAMES
f93b087d7f6d   ubuntu:latest  "/bin/bash"             39 minutes ago Up 39 minutes                c2-jegan
5a8c2b4e4692   ubuntu:latest  "/bin/bash"             41 minutes ago Up 41 minutes                c1-jegan
jegan@tektutor.org:~$
jegan@tektutor.org:~$ docker container inspect c1-jegan
[
  {
    "Id": "5a8c2b4e46924008351d1d41c4bcd51544b0b262c69ae46573729ad094ca2f9f",
    "Created": "2026-06-29T10:31:37.081139843Z",
    "Path": "/bin/bash",
    "Args": [],
    "State": {
      "Status": "running",
      "Running": true,
      "Paused": false,
      "Restarting": false,
      "OOMKilled": false,
      "Dead": false,
      "Pid": 53001,
      "ExitCode": 0,
      "Error": "",
      "StartedAt": "2026-06-29T10:31:37.116359236Z",
      "FinishedAt": "0001-01-01T00:00:00Z"
    },
    "Image": "sha256:5a4c6b929c57abe310fad22db2820f1423e645cdc9344bb05adaca9b50c3403f",
    "ResolvConfPath": "/var/lib/docker/containers/5a8c2b4e46924008351d1d41c4bcd51544b0b262c69ae46573729ad094ca2f9f/resolv.conf",
    "HostnamePath": "/var/lib/docker/containers/5a8c2b4e46924008351d1d41c4bcd51544b0b262c69ae46573729ad094ca2f9f/hostname",
    "HostsPath": "/var/lib/docker/containers/5a8c2b4e46924008351d1d41c4bcd51544b0b262c69ae46573729ad094ca2f9f/hosts"
  }
]
```

```
Mon Jun 29 4:43 PM
jegan@tektutor:~
jegan@tektutor:~
"org.opencontainers.image.version": "26.04"
}
},
"NetworkSettings": {
  "SandboxID": "e451ed67dfc919c041ec85159dec162814022b6125accb5e014607192133c5fe",
  "SandboxKey": "/var/run/docker/netns/e451ed67dfc9",
  "Ports": {},
  "Networks": {
    "bridge": {
      "IPAMConfig": null,
      "Links": null,
      "Aliases": null,
      "DriverOpts": null,
      "GwPriority": 0,
      "NetworkID": "4e636f92ca261130ad5d421a9e5c7e8f8da49905a8b9482754d831e7103d7119",
      "EndpointID": "efea4f0d870a331996cdc9090f82eb3696b3cc693429e61485aec2a96e263bd9",
      "Gateway": "172.17.0.1",
      "IPAddress": "172.17.0.2",
      "MacAddress": "6a:66:ac:7f:54:1c",
      "IPPrefixLen": 16,
      "IPv6Gateway": "",
      "GlobalIPv6Address": "",
      "GlobalIPv6PrefixLen": 0,
      "DNSNames": null
    }
  }
}
}
jegan@tektutor.org:~$
```


Lab - Finding IP address of a container

```
docker inspect c1-jegan | grep IPA
docker inspect -f {{.NetworkSettings.IPAddress}} c1-jegan
docker exec -it c1-jegan hostname -i
```

```
Mon Jun 29 4:45 PM
jegan@tektutor:~
jegan@tektutor:~
"SandboxID": "e451ed67dfc919c041ec85159dec162814022b6125accb5e014607192133c5fe",
"SandboxKey": "/var/run/docker/netns/e451ed67dfc9",
"Ports": {},
"Networks": {
  "bridge": {
    "IPAMConfig": null,
    "Links": null,
    "Aliases": null,
    "DriverOpts": null,
    "GwPriority": 0,
    "NetworkID": "4e636f92ca261130ad5d421a9e5c7e8f8da49905a8b9482754d831e7103d7119",
    "EndpointID": "efea4f0d870a331996cdc9090f82eb3696b3cc693429e61485aec2a96e263bd9",
    "Gateway": "172.17.0.1",
    "IPAddress": "172.17.0.2",
    "MacAddress": "6a:66:ac:7f:54:1c",
    "IPPrefixLen": 16,
    "IPv6Gateway": "",
    "GlobalIPv6Address": "",
    "GlobalIPv6PrefixLen": 0,
    "DNSNames": null
  }
}
}
jegan@tektutor.org:~$ docker inspect -f {{.NetworkSettings.Networks.bridge.IPAddress}} c1-jegan
172.17.0.2
jegan@tektutor.org:~$ docker inspect -f {{.NetworkSettings.SandboxKey}} c1-jegan
/var/run/docker/netns/e451ed67dfc9
jegan@tektutor.org:~$
```

Lab - Setup a Load balancer using containers

Docker - Port Forwarding



Let's create 3 web server containers using nginx:latest image from Docker Hub website

```
docker run -d --name nginx1-jegan --hostname nginx1-jegan nginx:latest
docker run -d --name nginx2-jegan --hostname nginx2-jegan nginx:latest
docker run -d --name nginx3-jegan --hostname nginx3-jegan nginx:latest
```

List all running containers

```
docker ps
```

Find the IP addresses of the 3 webserver containers

```
docker inspect nginx1-jegan | grep IPA
docker inspect nginx2-jegan | grep IPA
docker inspect nginx3-jegan | grep IPA
```

See if you are able to access the web pages from those web server containers

```
curl http://172.17.0.2:80
curl http://172.17.0.3:80
curl http://172.17.0.4:80
```

```
Mon Jun 29 5:04 PM
jegan@tektutor: ~
jegan@tektutor: ~
jegan@tektutor.org:~$ docker run -d --name nginx1-jegan --hostname nginx1-jegan nginx:latest
8d1751568faa2cb0c557d06271334a5a38632b2029d779c7b3db23ac1fa518b0
jegan@tektutor.org:~$
jegan@tektutor.org:~$ docker run -d --name nginx2-jegan --hostname nginx2-jegan nginx:latest
86f7557ff0f3dfb27d39a741dc29053e96809a910ae43d39bd6cc7fbc05b340
jegan@tektutor.org:~$
jegan@tektutor.org:~$ docker run -d --name nginx3-jegan --hostname nginx3-jegan nginx:latest
f43aac8a663087ed41edfaede99ee6d92f74f10c9a9ced6933ce208a9a13330c
jegan@tektutor.org:~$
jegan@tektutor.org:~$ docker ps
CONTAINER ID   IMAGE      COMMAND                  CREATED        STATUS        PORTS        NAMES
f43aac8a6630   nginx:latest   "/docker-entrypoint..."  4 seconds ago   Up 3 seconds   80/tcp       nginx3-jegan
86f7557ff0f3   nginx:latest   "/docker-entrypoint..."  15 seconds ago   Up 14 seconds   80/tcp       nginx2-jegan
8d1751568faa   nginx:latest   "/docker-entrypoint..."  23 seconds ago   Up 23 seconds   80/tcp       nginx1-jegan
jegan@tektutor.org:~$
jegan@tektutor.org:~$ docker inspect nginx1-jegan | grep IPA
      "IPAMConfig": null,
      "IPAddress": "172.17.0.2",
jegan@tektutor.org:~$ docker inspect nginx2-jegan | grep IPA
      "IPAMConfig": null,
      "IPAddress": "172.17.0.3",
jegan@tektutor.org:~$ docker inspect nginx3-jegan | grep IPA
      "IPAMConfig": null,
      "IPAddress": "172.17.0.4",
jegan@tektutor.org:~$
jegan@tektutor.org:~$ curl http://172.17.0.2:80
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
```

```
Mon Jun 29 5:08 PM
jegan@tektutor: ~
jegan@tektutor: ~
jegan@tektutor.org:~$ docker inspect nginx3-jegan | grep IPA
      "IPAMConfig": null,
      "IPAddress": "172.17.0.4",
jegan@tektutor.org:~$
jegan@tektutor.org:~$ curl http://172.17.0.2:80
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
html { color-scheme: light dark; }
body { width: 35em; margin: 0 auto;
font-family: Tahoma, Verdana, Arial, sans-serif; }
</style>
</head>
<body>
<h1>Welcome to nginx!</h1>
<p>If you see this page, nginx is successfully installed and working.
Further configuration is required for the web server, reverse proxy,
API gateway, load balancer, content cache, or other features.</p>

<p>For online documentation and support please refer to
<a href="https://nginx.org/">nginx.org</a>.<br/>
To engage with the community please visit
<a href="https://community.nginx.org/">community.nginx.org</a>.<br/>
For enterprise grade support, professional services, additional
security features and capabilities please refer to
<a href="https://f5.com/nginx">f5.com/nginx</a>.</p>

<p><em>Thank you for using nginx.</em></p>
</body>
```

```
Mon Jun 29 5:08 PM
jegan@tektutor:~
jegan@tektutor:~$ curl http://172.17.0.4:80
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
html { color-scheme: light dark; }
body { width: 35em; margin: 0 auto;
font-family: Tahoma, Verdana, Arial, sans-serif; }
</style>
</head>
<body>
<h1>Welcome to nginx!</h1>
<p>If you see this page, nginx is successfully installed and working.
Further configuration is required for the web server, reverse proxy,
API gateway, load balancer, content cache, or other features.</p>

<p>For online documentation and support please refer to
<a href="https://nginx.org/">nginx.org</a>.<br/>
To engage with the community please visit
<a href="https://community.nginx.org/">community.nginx.org</a>.<br/>
For enterprise grade support, professional services, additional
security features and capabilities please refer to
<a href="https://f5.com/nginx">f5.com/nginx</a>.</p>

<p><em>Thank you for using nginx.</em></p>
</body>
</html>
jegan@tektutor.org:~$ <!DOCTYPE html>
```

Now, let's create the load balancer container with port-forward(to make it accessible outside the system where the lb container is running)

```
docker run -d --name lb-jegan --hostname lb-jegan -p 8080:80 nginx:latest
```

In case you wish docker to identify an available port on your lab machine and forward to that port

```
docker rm -f lb-jegan
docker run -d --name lb-jegan --hostname lb-jegan -P nginx:latest
```

```
Mon Jun 29 5:41 PM
jegan@tektutor:~
jegan@tektutor:~
jegan@tektutor.org:~$ docker run -d --name lb-jegan --hostname lb-jegan -P nginx:latest
a40da94fab0d9374e845493733a4c5578a064b7581220db10e3c3e5d5646dfdf
jegan@tektutor.org:~$
jegan@tektutor.org:~$ docker ps
CONTAINER ID   IMAGE     COMMAND                  CREATED        STATUS        PORTS                               NAMES
a40da94fab0d   nginx:latest  "/docker-entrypoint..."  2 seconds ago  Up 2 seconds  0.0.0.0:32769->80/tcp, [::]:32769->80/tcp  lb-jegan
f43aac8a6630   nginx:latest  "/docker-entrypoint..."  35 minutes ago  Up 35 minutes  80/tcp                                nginx3-jegan
86f7557ff0f3   nginx:latest  "/docker-entrypoint..."  35 minutes ago  Up 35 minutes  80/tcp                                nginx2-jegan
8d1751568faa   nginx:latest  "/docker-entrypoint..."  36 minutes ago  Up 36 minutes  80/tcp                                nginx1-jegan
jegan@tektutor.org:~$
```

We need to copy the nginx.conf file from the lb-jegan container to configure it work like a load balancer

```
docker cp lb-jegan:/etc/nginx/nginx.conf .
cat nginx.conf
```

```
Mon Jun 29 5:12 PM
jegan@tektutor:~ -- docker exec -it lb-jegan /bin/sh
jegan@tektutor:~
jegan@tektutor.org:~$ docker run -d --name lb-jegan --hostname lb-jegan -p 8080:80 nginx:latest
00b9192b12b8aa2e8266d8f371af5aac67f9e330ca12bb11f199d14f480abc4b
jegan@tektutor.org:~$
jegan@tektutor.org:~$ docker ps
CONTAINER ID   IMAGE     COMMAND                  CREATED        STATUS        PORTS                               NAMES
00b9192b12b8   nginx:latest  "/docker-entrypoint..."  2 seconds ago  Up 1 second   0.0.0.0:8080->80/tcp, [::]:8080->80/tcp  lb-jegan
f43aac8a6630   nginx:latest  "/docker-entrypoint..."  8 minutes ago  Up 8 minutes  80/tcp                                nginx3-jegan
86f7557ff0f3   nginx:latest  "/docker-entrypoint..."  8 minutes ago  Up 8 minutes  80/tcp                                nginx2-jegan
8d1751568faa   nginx:latest  "/docker-entrypoint..."  8 minutes ago  Up 8 minutes  80/tcp                                nginx1-jegan
jegan@tektutor.org:~$
jegan@tektutor.org:~$ # Let's find the nginx.conf inside the lb-jegan container
jegan@tektutor.org:~$ docker exec -it lb-jegan /bin/sh
# hostname
lb-jegan
# cd /etc/nginx
# ls
conf.d  fastcgi_params  mime.types  modules  nginx.conf  scgi_params  uwsgi_params
# cat nginx.conf

user nginx;
worker_processes auto;

error_log /var/log/nginx/error.log notice;
pid /run/nginx.pid;
```

```
Mon Jun 29 5:14 PM
jegan@tektutor:~
jegan@tektutor.org:~$ # Let's copy the nginx.conf from lb-jegan container from path /etc/nginx/nginx.conf to local machine
jegan@tektutor.org:~$ docker cp lb-jegan:/etc/nginx/nginx.conf .
Successfully copied 2.56kB to /home/jegan/.
jegan@tektutor.org:~$ cat nginx.conf

user nginx;
worker_processes auto;

error_log /var/log/nginx/error.log notice;
pid /run/nginx.pid;

events {
    worker_connections 1024;
}

http {
    include /etc/nginx/mime.types;
    default_type application/octet-stream;

    log_format main '$remote_addr - $remote_user [$time_local] "$request" '
        '$status $body_bytes_sent "$http_referer" '
        '"$http_user_agent" "$http_x_forwarded_for"';

    access_log /var/log/nginx/access.log main;

    sendfile on;
    #tcp_nopush on;

```

```
Mon Jun 29 5:15 PM
jegan@tektutor:~
jegan@tektutor.org:~$ cat nginx.conf

error_log /var/log/nginx/error.log notice;
pid /run/nginx.pid;

events {
    worker_connections 1024;
}

http {
    include /etc/nginx/mime.types;
    default_type application/octet-stream;

    log_format main '$remote_addr - $remote_user [$time_local] "$request" '
        '$status $body_bytes_sent "$http_referer" '
        '"$http_user_agent" "$http_x_forwarded_for"';

    access_log /var/log/nginx/access.log main;

    sendfile on;
    #tcp_nopush on;

    keepalive_timeout 65;

    #gzip on;

    include /etc/nginx/conf.d/*.conf;
}
jegan@tektutor.org:~$
```

You need find the IP addresses of your nginx web server containers and update the nginx.conf file as shown below


```
Mon Jun 29 5:19 PM
jegan@tektutor:~
jegan@tektutor.org:~$ cat nginx.conf
user nginx;
worker_processes auto;

error_log /var/log/nginx/error.log notice;
pid /run/nginx.pid;

events {
    worker_connections 1024;
}

http {
    upstream myapp1 {
        server 172.17.0.2:80;
        server 172.17.0.3:80;
        server 172.17.0.4:80;
    }

    server {
        listen 80;

        location / {
            proxy_pass http://myapp1;
        }
    }
}
jegan@tektutor.org:~$
```

You need to update the nginx.conf file with your web server IPs

```
user nginx;
worker_processes auto;

error_log /var/log/nginx/error.log notice;
pid /run/nginx.pid;

events {
    worker_connections 1024;
}

http {
    upstream myapp1 {
        server 172.17.0.2:80;
        server 172.17.0.3:80;
        server 172.17.0.4:80;
    }

    server {
        listen 80;

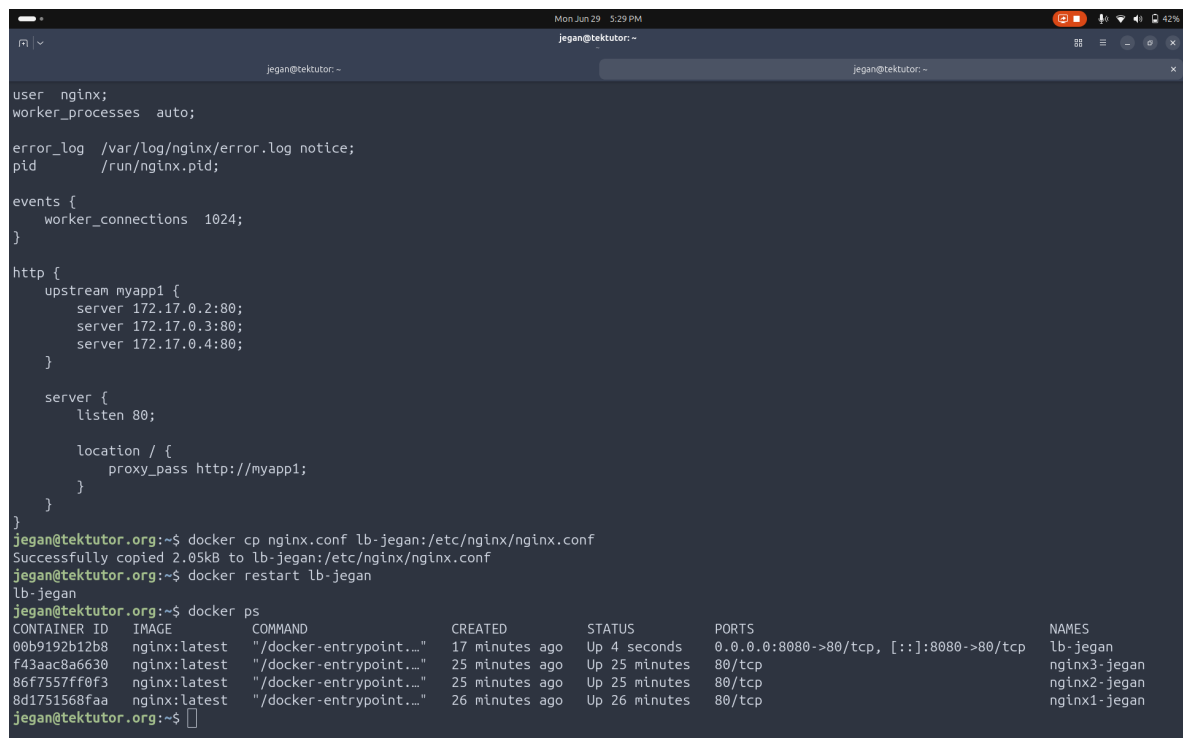
        location / {
            proxy_pass http://myapp1;
        }
    }
}
```

We need to copy this updated nginx.conf file from our local machine to the lb-jegan container

```
docker cp nginx.conf lb-jegan:/etc/nginx/nginx.conf
```

We need to restart the lb-jegan container to apply the config changes

```
docker restart lb-jegan
```



The screenshot shows a terminal window with the following content:

```
user nginx;
worker_processes auto;

error_log /var/log/nginx/error.log notice;
pid /run/nginx.pid;

events {
    worker_connections 1024;
}

http {
    upstream myapp1 {
        server 172.17.0.2:80;
        server 172.17.0.3:80;
        server 172.17.0.4:80;
    }

    server {
        listen 80;

        location / {
            proxy_pass http://myapp1;
        }
    }
}

jegan@tektutor.org:~$ docker cp nginx.conf lb-jegan:/etc/nginx/nginx.conf
Successfully copied 2.05kB to lb-jegan:/etc/nginx/nginx.conf
jegan@tektutor.org:~$ docker restart lb-jegan
lb-jegan
jegan@tektutor.org:~$ docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
00b9192b12b8	nginx:latest	"/docker-entrypoint..."	17 minutes ago	Up 4 seconds	0.0.0.0:8080->80/tcp, [::]:8080->80/tcp	lb-jegan
f43aac8a6630	nginx:latest	"/docker-entrypoint..."	25 minutes ago	Up 25 minutes	80/tcp	nginx3-jegan
86f7557ff0f3	nginx:latest	"/docker-entrypoint..."	25 minutes ago	Up 25 minutes	80/tcp	nginx2-jegan
8d1751568faa	nginx:latest	"/docker-entrypoint..."	26 minutes ago	Up 26 minutes	80/tcp	nginx1-jegan

Make sure the lb-jegan container is running after our config changes

```
docker ps
```


Let's customize the web page in nginx1-jegan, nginx2-jegan and nginx3-jegan web server containers

```
echo "<h1>Web Server 1</h1>" > index.html
docker cp index.html nginx1-jegan:/usr/share/nginx/html/index.html

echo "<h1>Web Server 2</h1>" > index.html
docker cp index.html nginx2-jegan:/usr/share/nginx/html/index.html

echo "<h1>Web Server 3</h1>" > index.html
docker cp index.html nginx3-jegan:/usr/share/nginx/html/index.html
```


```
curl http://172.17.0.2:80
curl http://172.17.0.3:80
curl http://172.17.0.4:80
```

A terminal window titled 'jegan@tektutor:~' showing the configuration of three nginx containers. The user creates an 'index.html' file with the content '<h1>Web Server 1</h1>' and copies it to the 'nginx1' container. This process is repeated for 'nginx2' and 'nginx3' with different server numbers. Finally, the user runs 'curl' commands to verify that each container is responding with its respective page title. The terminal output shows the file creation, copying process (2.05kB), and the successful curl responses for each container.


```
jegan@tektutor.org:~$ echo "<h1>Web Server 1</h1>" > index.html
jegan@tektutor.org:~$ cat index.html
<h1>Web Server 1</h1>
jegan@tektutor.org:~$ docker cp index.html nginx1-jegan:/usr/share/nginx/html/index.html
Successfully copied 2.05kB to nginx1-jegan:/usr/share/nginx/html/index.html
jegan@tektutor.org:~$
jegan@tektutor.org:~$ echo "<h1>Web Server 2</h1>" > index.html
jegan@tektutor.org:~$ docker cp index.html nginx2-jegan:/usr/share/nginx/html/index.html
Successfully copied 2.05kB to nginx2-jegan:/usr/share/nginx/html/index.html
jegan@tektutor.org:~$
jegan@tektutor.org:~$ echo "<h1>Web Server 3</h1>" > index.html
jegan@tektutor.org:~$ docker cp index.html nginx3-jegan:/usr/share/nginx/html/index.html
Successfully copied 2.05kB to nginx3-jegan:/usr/share/nginx/html/index.html
jegan@tektutor.org:~$
jegan@tektutor.org:~$ curl http://172.17.0.2:80
<h1>Web Server 1</h1>
jegan@tektutor.org:~$
jegan@tektutor.org:~$ curl http://172.17.0.3:80
<h1>Web Server 2</h1>
jegan@tektutor.org:~$ curl http://172.17.0.4:80
<h1>Web Server 3</h1>
jegan@tektutor.org:~$
```

Let's verify if the lb-jegan container is working as configured(as a load balancer) from firefox web browser on your lab machine


```
http://localhost:8080
http://localhost:8080
http://localhost:8080
```



Web Server 1



Web Server 2



Checking the load balancer logs

```
docker logs lb-jegan
```

```
jegan@tektutor.org:~$ docker ps
CONTAINER ID   IMAGE      COMMAND                  CREATED        STATUS        PORTS                               NAMES
00b9192b12b8   nginx:latest "/docker-entrypoint..." 23 minutes ago Up 5 minutes   0.0.0.0:8080->80/tcp, [::]:8080->80/tcp lb-jegan
f43aac8a6630   nginx:latest "/docker-entrypoint..." 31 minutes ago Up 31 minutes   80/tcp      nginx3-jegan
86f7557ff0f3   nginx:latest "/docker-entrypoint..." 31 minutes ago Up 31 minutes   80/tcp      nginx2-jegan
8d1751568faa   nginx:latest "/docker-entrypoint..." 31 minutes ago Up 31 minutes   80/tcp      nginx1-jegan

jegan@tektutor.org:~$ docker logs lb-jegan
/docker-entrypoint.sh: /docker-entrypoint.d/ is not empty, will attempt to perform configuration
/docker-entrypoint.sh: Looking for shell scripts in /docker-entrypoint.d/
/docker-entrypoint.sh: Launching /docker-entrypoint.d/10-listen-on-ipv6-by-default.sh
10-listen-on-ipv6-by-default.sh: info: Getting the checksum of /etc/nginx/conf.d/default.conf
10-listen-on-ipv6-by-default.sh: info: Enabled listen on IPv6 in /etc/nginx/conf.d/default.conf
/docker-entrypoint.sh: Sourcing /docker-entrypoint.d/15-local-resolvers.envsh
/docker-entrypoint.sh: Launching /docker-entrypoint.d/20-envsubst-on-templates.sh
/docker-entrypoint.sh: Launching /docker-entrypoint.d/30-tune-worker-processes.sh
/docker-entrypoint.sh: Configuration complete; ready for start up
2026/06/29 11:41:31 [notice] 1#1: using the "epoll" event method
2026/06/29 11:41:31 [notice] 1#1: nginx/1.31.1
2026/06/29 11:41:31 [notice] 1#1: built by gcc 14.2.0 (Debian 14.2.0-19)
2026/06/29 11:41:31 [notice] 1#1: OS: Linux 7.0.0-22-generic
2026/06/29 11:41:31 [notice] 1#1: getrlimit(RLIMIT_NOFILE): 1024:524288
2026/06/29 11:41:31 [notice] 1#1: start worker processes
2026/06/29 11:41:31 [notice] 1#1: start worker process 29
2026/06/29 11:41:31 [notice] 1#1: start worker process 30
2026/06/29 11:41:31 [notice] 1#1: start worker process 31
2026/06/29 11:41:31 [notice] 1#1: start worker process 32
2026/06/29 11:41:31 [notice] 1#1: start worker process 33
2026/06/29 11:41:31 [notice] 1#1: start worker process 34
2026/06/29 11:41:31 [notice] 1#1: start worker process 35
2026/06/29 11:41:31 [notice] 1#1: start worker process 36
2026/06/29 11:41:31 [notice] 1#1: start worker process 37
2026/06/29 11:41:31 [notice] 1#1: start worker process 38
2026/06/29 11:41:31 [notice] 1#1: start worker process 39
2026/06/29 11:41:31 [notice] 1#1: start worker process 40
2026/06/29 11:41:31 [notice] 1#1: start worker process 41
2026/06/29 11:41:31 [notice] 1#1: start worker process 42
```

Checking the webserver logs

```
docker logs nginx1-jegan
docker logs nginx2-jegan
docker logs nginx3-jegan
```

```

jegan@tektutor:~$ docker ps
CONTAINER ID   IMAGE          COMMAND                  CREATED        STATUS        PORTS                               NAMES
00b9192b12b8   nginx:latest   "/docker-entrypoint..." 24 minutes ago Up 6 minutes   0.0.0.0:8080->80/tcp, [::]:8080->80/tcp   lb-jegan
f43aac8a6630   nginx:latest   "/docker-entrypoint..." 32 minutes ago Up 32 minutes   80/tcp                                   nginx3-jegan
86f7557ff0f3   nginx:latest   "/docker-entrypoint..." 32 minutes ago Up 32 minutes   80/tcp                                   nginx2-jegan
8d1751568faa   nginx:latest   "/docker-entrypoint..." 32 minutes ago Up 32 minutes   80/tcp                                   nginx1-jegan

jegan@tektutor.org:~$ docker logs nginx1-jegan
/docker-entrypoint.sh: /docker-entrypoint.d/ is not empty, will attempt to perform configuration
/docker-entrypoint.sh: Looking for shell scripts in /docker-entrypoint.d/
/docker-entrypoint.sh: Launching /docker-entrypoint.d/10-listen-on-ipv6-by-default.sh
10-listen-on-ipv6-by-default.sh: info: Getting the checksum of /etc/nginx/conf.d/default.conf
10-listen-on-ipv6-by-default.sh: info: Enabled listen on IPv6 in /etc/nginx/conf.d/default.conf
/docker-entrypoint.sh: Sourcing /docker-entrypoint.d/15-local-resolvers.envsh
/docker-entrypoint.sh: Launching /docker-entrypoint.d/20-envsubst-on-templates.sh
/docker-entrypoint.sh: Launching /docker-entrypoint.d/30-tune-worker-processes.sh
/docker-entrypoint.sh: Configuration complete; ready for start up
2026/06/29 11:33:00 [notice] 1#1: using the "epoll" event method
2026/06/29 11:33:00 [notice] 1#1: nginx/1.31.1
2026/06/29 11:33:00 [notice] 1#1: built by gcc 14.2.0 (Debian 14.2.0-19)
2026/06/29 11:33:00 [notice] 1#1: OS: Linux 7.0.0-22-generic
2026/06/29 11:33:00 [notice] 1#1: getrlimit(RLIMIT_NOFILE): 1024:524288
2026/06/29 11:33:00 [notice] 1#1: start worker processes
2026/06/29 11:33:00 [notice] 1#1: start worker process 29
2026/06/29 11:33:00 [notice] 1#1: start worker process 30
2026/06/29 11:33:00 [notice] 1#1: start worker process 31

```

Lab - Create a mysql db container

```
docker run -d --name mysql-jegan --hostname mysql-jegan -e
MYSQL_ROOT_PASSWORD=root@123 mysql:latest
docker ps
docker logs mysql-jegan
```

```

nt script for MySQL Server 9.7.0-1.el
izing database files
ver] MySQL Server Initialization - st
ver] /usr/sbin/mysqld (mysqld 9.7.0)
oDB] InnoDB initialization has starte
oDB] InnoDB initialization has ended.
rver] root@localhost is created with
ver] MySQL Server Initialization - en
files initialized

```

```
Tue Jun 30 10:54 AM
jegan@tektutor: ~
jegan@tektutor: ~
jegan@tektutor: ~
2026-06-30T05:21:43.788053Z 0 [System] [MY-010931] [Server] /usr/sbin/mysqld: ready for connections. Version: '9.7.0' socket: '/var/run/mysqld/mysqld.sock' port: 0 MySQL Community Server - GPL.
2026-06-30T05:21:43.788069Z 0 [System] [MY-011323] [Server] X Plugin ready for connections. Socket: /var/run/mysqld/mysqdx.sock
2026-06-30 05:21:43+00:00 [Note] [Entrypoint]: Temporary server started.
'/var/lib/mysql/mysql.sock' -> '/var/run/mysqld/mysqld.sock'
Warning: Unable to load '/usr/share/zoneinfo/iso3166.tab' as time zone. Skipping it.
Warning: Unable to load '/usr/share/zoneinfo/leap-seconds.list' as time zone. Skipping it.
Warning: Unable to load '/usr/share/zoneinfo/leapseconds' as time zone. Skipping it.
Warning: Unable to load '/usr/share/zoneinfo/tzdata.zi' as time zone. Skipping it.
Warning: Unable to load '/usr/share/zoneinfo/zone.tab' as time zone. Skipping it.
Warning: Unable to load '/usr/share/zoneinfo/zone1970.tab' as time zone. Skipping it.

2026-06-30 05:21:44+00:00 [Note] [Entrypoint]: Stopping temporary server
2026-06-30T05:21:44.334183Z 11 [System] [MY-013172] [Server] Received SHUTDOWN from user root. Shutting down mysqld (Version: 9.7.0).
2026-06-30T05:21:45.269926Z 0 [System] [MY-010910] [Server] /usr/sbin/mysqld: Shutdown complete (mysqld 9.7.0) MySQL Community Server - GPL.
2026-06-30T05:21:45.269933Z 0 [System] [MY-015016] [Server] MySQL Server - end.
2026-06-30 05:21:45+00:00 [Note] [Entrypoint]: Temporary server stopped

2026-06-30 05:21:45+00:00 [Note] [Entrypoint]: MySQL init process done. Ready for start up.

2026-06-30T05:21:45.350470Z 0 [System] [MY-015015] [Server] MySQL Server - start.
2026-06-30T05:21:45.542335Z 0 [System] [MY-010116] [Server] /usr/sbin/mysqld (mysqld 9.7.0) starting as process 1
2026-06-30T05:21:45.542344Z 0 [System] [MY-015603] [Server] MySQL Server has access to 24 logical CPUs.
2026-06-30T05:21:45.542352Z 0 [System] [MY-015603] [Server] MySQL Server has access to 65746083840 bytes of physical memory.
2026-06-30T05:21:45.552808Z 1 [System] [MY-013576] [InnoDB] InnoDB initialization has started.
2026-06-30T05:21:45.675170Z 1 [System] [MY-013577] [InnoDB] InnoDB initialization has ended.
2026-06-30T05:21:45.814419Z 0 [Warning] [MY-010068] [Server] CA certificate ca.pem is self signed.
2026-06-30T05:21:45.814432Z 0 [System] [MY-013602] [Server] Channel mysql_main configured to support TLS. Encrypted connections are now supported for this channel.
2026-06-30T05:21:45.815677Z 0 [Warning] [MY-011810] [Server] Insecure configuration for --pid-file: Location '/var/run/mysqld' in the path is accessible to all OS users. Consider choosing a different directory.
2026-06-30T05:21:45.827128Z 0 [System] [MY-010931] [Server] /usr/sbin/mysqld: ready for connections. Version: '9.7.0' socket: '/var/run/mysqld/mysqld.sock' port: 3306 MySQL Community Server - GPL.
2026-06-30T05:21:45.827136Z 0 [System] [MY-011323] [Server] X Plugin ready for connections. Bind-address: '::' port: 33060, socket: /var/run/mysqld/mysqdx.sock
jegan@tektutor.org:~$
```

Let's get inside the mysql db server container

```
docker exec -it mysql-jegan /bin/sh
mysql -u root -p
SHOW DATABASES;
CREATE DATABASE tektutor;
USE tektutor;
CREATE TABLE trainings ( id INT NOT NULL, name VARCHAR(200) NOT NULL, duration
VARCHAR(200) NOT NULL, PRIMARY KEY(id) );
INSERT INTO trainings VALUES ( 1, "DevOps", "5 Days" );
INSERT INTO trainings VALUES ( 2, "Microservices with Golang", "5 Days" );
SELECT * FROM trainings;
exit
exit
docker ps
```

```
Tue Jun 30 10:57 AM
jegan@tektutor:~
jegan@tektutor.org:~$ # Let's get inside the mysql db server container
jegan@tektutor.org:~$ docker exec -it mysql-jegan /bin/sh
sh-5.1#
sh-5.1# hostname
sh: hostname: command not found
sh-5.1# mysql -u root -p
Enter password:
Welcome to the MySQL monitor.  Commands end with ; or \g.
Your MySQL connection id is 9
Server version: 9.7.0 MySQL Community Server - GPL

Copyright (c) 2000, 2026, Oracle and/or its affiliates.

Oracle is a registered trademark of Oracle Corporation and/or its
affiliates. Other names may be trademarks of their respective
owners.

Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.

mysql> SHOW DATABASES;
+-----+
| Database |
+-----+
| information_schema |
| mysql |
| performance_schema |
| sys |
+-----+
4 rows in set (0.002 sec)

mysql> CREATE DATABASE tektutor;
Query OK, 1 row affected (0.011 sec)

mysql> U
->
-> ;
ERROR 1064 (42000): You have an error in your SQL syntax; check the manual that corresponds to your MySQL server version for the right syntax to u
```

```
Tue Jun 30 10:58 AM
jegan@tektutor:~
jegan@tektutor.org:~$
->
-> ;
ERROR 1064 (42000): You have an error in your SQL syntax; check the manual that corresponds to your MySQL server version for the right syntax to u
se near 'U' at line 1
mysql> USE tektutor;
Database changed
mysql> SHOW t
-> ;
ERROR 1064 (42000): You have an error in your SQL syntax; check the manual that corresponds to your MySQL server version for the right syntax to u
se near 't' at line 1
mysql> SHOW TABLES;
Empty set (0.001 sec)

mysql> CREATE TABLE trainings ( id INT NOT NULL, name VARCHAR(200) NOT NULL, duration VARCHAR(200) NOT NULL, PRIMARY KEY(id) );
Query OK, 0 rows affected (0.018 sec)

mysql> INSERT INTO trainings VALUES ( 1, "DevOps", "5 Days" );
Query OK, 1 row affected (0.013 sec)

mysql> INSERT INTO trainings VALUES ( 2, "Microservices with Golang", "5 Days" );
Query OK, 1 row affected (0.011 sec)

mysql> SELECT * FROM trainings;
+-----+
| id | name | duration |
+-----+
| 1 | DevOps | 5 Days |
| 2 | Microservices with Golang | 5 Days |
+-----+
2 rows in set (0.000 sec)

mysql> exit
Bye
sh-5.1# exit
exit
jegan@tektutor.org:~$
```

Let's try to restart the mysql container

```
docker restart mysql-jegan
docker ps
docker exec -it mysql-jegan /bin/sh
mysql -u root
SHOW DATABASES;
```



```
USE tektutor;  
SHOW TABLES;  
SELECT * FROM trainings;  
exit  
exit
```

```

jegan@tektutor:~$ docker ps
CONTAINER ID   IMAGE          COMMAND                  CREATED        STATUS        PORTS                    NAMES
ab4009975448   mysql:latest   "docker-entrypoint.s..." 7 minutes ago  Up 7 minutes  3306/tcp, 33060/tcp      mysql-jegan
jegan@tektutor:~$ docker restart mysql-jegan
mysql-jegan
jegan@tektutor:~$ docker ps
CONTAINER ID   IMAGE          COMMAND                  CREATED        STATUS        PORTS                    NAMES
ab4009975448   mysql:latest   "docker-entrypoint.s..." 9 minutes ago  Up 4 seconds  3306/tcp, 33060/tcp      mysql-jegan
jegan@tektutor:~$ docker exec -it mysql-jegan /bin/sh
sh-5.1# mysql -u root -p
Enter password:
Welcome to the MySQL monitor.  Commands end with ; or \g.
Your MySQL connection id is 9
Server version: 9.7.0 MySQL Community Server - GPL

Copyright (c) 2000, 2026, Oracle and/or its affiliates.

Oracle is a registered trademark of Oracle Corporation and/or its
affiliates. Other names may be trademarks of their respective
owners.

Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.

mysql> SHOW DATABASES;
+-----+
| Database |
+-----+
| information_schema |
| mysql |
| performance_schema |
| sys |
| tektutor |
+-----+
5 rows in set (0.009 sec)

mysql> USE tektutor;
Reading table information for completion of table and column names

```

```

+-----+
| information_schema |
| mysql              |
| performance_schema |
| sys                |
| tektutor            |
+-----+
5 rows in set (0.009 sec)

mysql> USE tektutor;
Reading table information for completion of table and column names
You can turn off this feature to get a quicker startup with -A

Database changed
mysql> SHOW TABLES;
+-----+
| Tables_in_tektutor |
+-----+
| trainings           |
+-----+
1 row in set (0.001 sec)

mysql> SELECT * FROM trainings;
+-----+
| id | name                                | duration |
+-----+
| 1  | DevOps                             | 5 Days  |
| 2  | Microservices with Golang          | 5 Days  |
+-----+
2 rows in set (0.000 sec)

mysql> exit
Bye
sh-5.1# exit
exit
jegan@tektutor.org:~$

```

Let's delete the mysql container. At this point, we not only lost the container, we also lost the data stored inside the container.

```
docker rm -f mysql-jegan
```

Hence, we must always store application data in an external storage.

```
mkdir -p /tmp/jegan/mysql
docker run -d --name mysql-jegan --hostname mysql-jegan -e
MYSQL_ROOT_PASSWORD=root@123 -v /tmp/jegan/mysql:/var/lib/mysql mysql:latest
docker exec -it mysql-jegan /bin/sh
mysql -u root -p
SHOW DATABASES;
CREATE DATABASE tektutor;
USE tektutor;
CREATE TABLE trainings ( id INT NOT NULL, name VARCHAR(200) NOT NULL, duration
VARCHAR(200) NOT NULL, PRIMARY KEY(id) );
INSERT INTO trainings VALUES ( 1, "DevOps", "5 Days" );
INSERT INTO trainings VALUES ( 2, "Microservices with Golang", "5 Days" );
SELECT * FROM trainings;
exit
exit
docker ps
```

Let's delete the mysql-jegan container and recreate a new one

```
docker rm mysql-jegan
docker run -d --name mysql-jegan --hostname mysql-jegan -e
MYSQL_ROOT_PASSWORD=root@123 -v /tmp/jegan/mysql:/var/lib/mysql mysql:latest
docker exec -it mysql-jegan /bin/sh
mysql -u root -p
SHOW DATABASES;
USE tektutor;
SHOW TABLES;
SELECT * FROM trainings;
exit
exit
```

Though we delete the container and recreated a new one, we didn't lose the data because we stored the database in an external disk.
This is how, containers are used in real-world applications.

Day 2

Info - Container Orchestration Platform

- in order to deploy your application into Container Orchestration Platform, we must first containerize them
- In real world applications, no one manages containerized application workloads manually
- In real world, Container Orchestration Platforms manage our containerized application workloads
- They provide
 - in-built monitoring, health-check, readiness check, liveness check for your application
 - self-healing to your applications
 - a way to scale up/down your application instances based on user traffic
 - rolling update
 - upgrading your application from one version to other without any downtime
 - rolling back to earlier stable versions are also possible
 - ways to collect performance metrics
- any container orchestration platform works as a cluster of many nodes
- nodes are nothing but Physical Servers or Virtual Machines in private/public cloud
- examples
 - Docker SWARM
 - Kubernetes
 - Rancher
 - Red Hat Openshift
 - AWS eks
 - AWS ROSA
 - Azure aks
 - Azure ARO

Info - Docker SWARM

- is developed by Docker Inc as an open-source project
- it supported only Docker Containerized applications
- it is a very light-weight setup that works in laptop/dekstop with


normal configurations

- it is ideal setup for learning purpose
- it is ideal for Dev/QA environment
- it is not production-grade


Info - Kubernetes


- is a Container Orchestration platform developed in Golang by Google
- initially it was called borg, Google has been using this container orchestration platforms internally within google for several years
- it is production-grade because google used it on complex projects many years
- it supports many different container runtimes/engines
- Kubernetes supports something called CRI(Container Runtime Interface)
- any container runtime/engine that must be supported by Kubernetes they have to implement the CRI
- Kubernetes can/will interact with any type of Container Runtime/Engine via CRI
- Kubernetes supports 2 types of nodes
 - Master Node
 - this is where Control Plane components will be running
 - Control Plane Components
 1. API Server
 2. etcd key/value database
 3. scheduler
 4. Controller Managers(collection of many controllers)
 - Worker Node
 - this is where user applications will be running
- Kubernetes is primarily a CLI application
- Kubernetes doesn't support
 - a production-grade webconsole
 - User Management
 - Internal Container Registry (out of box)
 - Prometheus/Grafana performance metrics collection (out of box)
- Kubernetes supports basic building blocks to extend the Kubernetes features by adding
 - your own custom resources
 - your own custom controllers
- Kubernetes allows one to deploy their application only from a Container Image
 - In Kubernetes Master/Worker nodes, we can choose to install any Linux OS distribution (ubuntu, rocky, fedora, RHEL, etc.,)

Info - Kubernetes High-Level Architecture



Info - Red Hat Openshift High-Level Architecture





Info - Red Hat Openshift

- is a Red Hat's distribution of Kubernetes
- Openshift is developed on top of opensource Kubernetes
- Hence, Openshift is a superset of Kubernetes
- Openshift is Kubernetes with batteries included
- Openshift comes with these additional features on top of what is already supported by Kubernetes
 1. User Management
 2. Production Web-console (GUI)
 3. Many additional features
 - Route
 - S2I (Source to Image)
 - one can deploy their application from plain source code coming from GitHub or any other version control
 - supports many different strategies
 - docker
 - source
 - pipeline
 4. Internal Container Registry
- Red Hat acquired a company named CoreOS
 - CoreOS organization had several interesting products
 - rkt(pronounced as Rocket) - container runtime

- CoreOS - Linux based minimal & secure operating system good enough for Container Orchestration Platform
- Network addons - Flannel
- Red Hat Openshift starting from v4.x,
 - supports only Red Hat Enterprise CoreOS operating (RHCOS) in Master Nodes
 - supports either Red Hat Enterprise Linux(RHEL) or RHCOS in Worker Node
 - Red Hat Openshift generally recommends to install RHCOS in all nodes
 - supports only CRI-O Container Runtime
 - RHCOS Operating System
- Create
- Describe

Finding IP Address of a Pod

- is an Container Orchestration Platform optimized minimal and secured Operating System based on RHEL
- this OS comes with CRI-O & Podman pre-installed
- it is an immutable OS many restrictions
- some of the folders like /var, /bin, /usr/bin are all treated as read-only folder
- only through Machine Config Operators we can update certain reserved folders, meaning application won't be able to modify those restricted folders
- Ports 0-1024 are reserved for Openshift's internal use

Info - Pod

- Pod is a logical grouping of related containers
- Pod is record stored in etcd database
 - is a configuration object stored in etcd database
 - what really runs on nodes is containers
- Unlike Docker container where each container get a dedicated IP address, in case of Pod, all containers that are part of a Pod they all share the same Private IP address
- applications runs with a Pod container
- Every Pod get its own port range i.e 0 to 65535
- if there are two container c1 and c2 within Pod P1, if c1 is using port 8080 then c2 won't be able to use port 8080
- the containers within the same Pod, they can communicate with each with regular IPC mechanisms

Info - Types of applications that can be deployed into Openshift

- Stateless applications
 - k8s/Openshift supports a resource called Deployment
- Stateful applications
 - k8s/Openshift supports a resource called StatefulSet
- Application that performs a one-time activity and stops in some time
 - Job
- Application that must be invoked periodically on a particular day and time
 - CronJob (this internally used Job)
- Applications that run one instance per Node
 - DaemonSet

Info - API Server

- is the heart of kubernetes/openshift
- anything and everything in Kubernetes/Openshift is co-ordinated by API server, hence this is one of the most critical components
- API server only performs database activities like
 - it creates a new record in etcd datastore
 - it updates an existing record in the etcd datastore
 - it retrieves an existing record from the etcd datastore
 - it deletes an existing record from the etcd datastore
 - everytime some new record are added, existing records updated, existing records deleted it will send events to registered controllers

Info - etcd

- it is distributed database which stores data in key/value format
- it is an idendepent opensource project used by Kubernetes & openshift
- this can be used outside the scope of K8s/Openshift as well

Info - Scheduler

- is a Pod
- is one of the Control Plane components that runs in the master node

- it is responsible to identify a healthy node where a new Pod can be deployed
- scheduler by itself can't deploy a pod into any node, it is allowed only to share it scheduling recommendations to API Server via REST call

Info - Controller Managers

- the real doers are controllers
- is a collection of many in-built controllers
- for each type of resource supported in Openshift/Kubernetes there is one Controller
- For example
 - Deployment Controller manages Deployment resource (stateless application)
 - ReplicaSet Controller manages ReplicaSet resource
 - Job Controller manages Job resource
 - StatefulSet Controller manages StatefulSet resource (stateful application)
 - DaemonSet Controller manages DaemonSet resource
- Controller is a application that runs within a Pod container
- they have unrestricted access within the Kubernetes/Openshift cluster
- they can monitor specific type of resources created/updated/deleted under any project within the cluster
- they get notified by API Server via events
- Controllers if it needs to communicate something to API Server, it make a REST API call

Lab - Login to openshift from command-line

```
cat ~/openshift.txt
oc login --username=kubeadmin --password=VVnz6-ezZ9W-RVbw4-NPK3R https://
api.ocp4.palmeto.org:6443 --insecure-skip-tls-verify
```

Lab - Listing Nodes in Red Hat Openshift

```
oc version
kubectl version

oc get nodes
kubectl get nodes
```

```
oc get nodes -o wide
kubectl get nodes -o wide
```

Lab - Finding additional meta-data about a node

```
oc describe node/master01.ocp4.palmeto.org
oc describe node/worker01.ocp4.palmeto.org
```

Lab - List and find all control plane pods

```
oc get pods --all-namespaces -o wide | grep api
oc get pods --all-namespaces -o wide | grep scheduler
```

Lab - Managing projects in Openshift

- each team in an organization setup, will create atleast 1 project to isolate their application deployment from the other teams in the same organization
- generally, a single Openshift cluster will be shared by all the team in an organization
- openshift, supports features to restrict access to only those team members working a project

Create a new-project

```
oc new-project jegan
```

List all projects

```
oc get projects
oc get project
```

Switching between projects

```
oc project default
oc project jegan
```

Finding the currently active project

```
oc project
```

Deleting a project(please delete only your project)

```
oc delete project jegan
```

Lab - Deploying your first application into Openshift using imperative approach

```
oc new-project jegan-project
oc create deploy nginx --image=image-registry.openshift-image-registry.svc:5000/
openshift/bitnami-nginx:1.26 --replicas=3

# Listing deployments
oc get deployments
oc get deployment
oc get deploy

# Listing replicaset
oc get replicaset
oc get rs

# Listing pods
oc get pods
oc get pod
oc get po

# Listing multiple resource with a single command
oc get deploy,rs,po
oc get all

# Finding more details about your deployment
oc describe deploy/nginx

# Finding more details about replicaset
oc describe rs/nginx-dbf56c96

# Finding more details about the pod
oc describe pod/nginx-dbf56c96-l296j

# Editing a deployment and update its replicas from 3 to 5, save and exit
oc edit deploy/nginx
```

```
oc get pods

oc get deploy/nginx -o yaml
oc get deploy/nginx -o json

oc get rs/nginx-dbf56c96 -o yaml
oc get pod/nginx-dbf56c96-jfx2m -o yaml
```

Lab - Scale up/down your deployments

```
oc project jegan-project
# Scale up
oc scale deploy/nginx --replicas=5

# Scale down
oc scale deploy/nginx --replicas=3
```

Lab - Port forward for quick testing/debugging (Not for production environment)

```
oc project jegan-project
oc get pods
# Terminal 1
oc port-forward pod/hello-7cf46f6476-fwzrn 9090:8080

# Terminal 2
curl http://localhost:9090

# To stop the port forward after your testing
# Press Ctrl+C in Terminal 1
# From Terminal 2, curl will no more work
```

Lab - Getting inside a node shell (only a cluster/kube admin do can do this at work place)

```
oc get nodes
oc debug node/master01.ocp4.palmeto.org
chroot /host
podman --version
crictl --version
crictl images
```

```
cricctl ps
exit
exit
```

Lab - Getting inside a pod shell (this doesn't require admin access)

```
oc project jegan-project
oc get pods
oc rsh pod/nginx-dbf56c96-wtbp7
exit

oc rsh deploy/nginx
exit
```

Lab - Creating an internal clusterip service from your nginx deployment

```
oc project jegan-project
oc get deploy -l app=nginx
oc expose deploy/nginx --type=ClusterIP --port=8080

oc get services
oc get service
oc get svc

oc get endpoints

oc describe svc/nginx

oc scale deploy/nginx --replicas=8

oc describe svc/nginx
oc get endpoints

oc scale deploy/nginx --replicas=3
```

```
oc describe svc/nginx
oc get endpoints
```

Day 3

Info - You can find Red Hat Base images here

```
https://catalog.redhat.com/en/software/containers/explore
```

Info - You can find other Red Hat Openshift compatible images here

```
https://catalog.redhat.com/en/search?searchType=Containers
```


Info - What happens internally in Openshift when we deploy an application

```
oc project jegan-project
oc create deploy nginx --image=image-registry.openshift-image-registry.svc:5000/
openshift/bitnami-nginx:1.26 --replicas=3
```

Note

- oc client tool makes a REST call to API Server requesting the API Server to create a deployment
- Once API Server receives the request from oc client, it creates a Deployment database entry in etcd database
- API Server then sends broadcasting event saying new Deployment created along with deployment details
- Deployment Controller receives this event, it then makes a REST call to API Server requesting it to create a ReplicaSet for the nginx deployment
- API Server creates a ReplicaSet db entry(new record) in the etcd database
- API Server sends a broadcast event saying new ReplicaSet created
- ReplicaSet Controller receives the event, it then makes a REST call to API Server requesting it to create 3 Pods
- API Server create 3 Pod records in the etcd database

- API Server sends broadcast event for each new Pod created in the etcd database
- Scheduler receives the event, it then identifies a healthy node where the new Pod can be deployed
- Scheduler makes a REST call to API Server to send it scheduling recommendation. This will be done for each Pod.
- API Server receives the scheduling recommendations from Scheduler, it then retrieves the Pod record from etcd and updates its status as Scheduled to so and so node
- API Server sends a broadcasting event saying Pod1 scheduled to Worker01 node, this happens for each Pod.
- Kubelet Container Agent that runs on Worker01 node receives the event, it then pulls the container image, creates and starts the container on Worker01
- Kubelet monitors the status of the Container created for Pod1, and it periodically updates the status back to API Server in a heart-beat fashion
- API Server receives these updates, retrieves the Pod database entry from etcd and updates the Pod status



Lab - Deploying a stateless application in declarative style

```
# Delete your existing project
oc delete project jegan

# Create a new project
oc new-project jegan

# Generate the declarative manifest file to deploy nginx
```



```

oc create deploy nginx --image=image-registry.openshift-image-registry.svc:5000/
openshift/bitnami-nginx:1.26 --replicas=3 --dry-run=client -o yaml

# Redirect the output shown to a yaml file
oc create deploy nginx --image=image-registry.openshift-image-registry.svc:5000/
openshift/bitnami-nginx:1.26 --replicas=3 --dry-run=client -o yaml > nginx-
deploy.yml

# Deploy the nginx in declarative fashion
oc create -f nginx-deploy.yml --save-config

# In the template section of the nginx-deploy.yml add additional labels and apply
oc apply -f nginx-deploy.yml

# List the deploy,replicaset and pods
oc get deploy,rs,po

## Delete the deployment declaratively
oc delete -f nginx-deploy.yml
oc get deploy,rs,po

```

Lab - Creating a ClusterIP Internal Service in declarative style

Create a pod.yml with code below

```

apiVersion: v1
kind: Pod
metadata:
  labels:
    app: test
    name: test-pod
spec:
  containers:
    - image: image-registry.openshift-image-registry.svc:5000/openshift/
      hello:latest
      imagePullPolicy: IfNotPresent
      name: test

```

Create the test pod

```

oc project jegan-project
oc apply -f pod.yml
oc get pods

```

Let's create the nginx deployment in declarative style

```
oc create deploy nginx --image=image-registry.openshift-image-registry.svc:5000/
openshift/bitnami-nginx:1.26 --replicas=3 --dry-run=client -o yaml
oc create deploy nginx --image=image-registry.openshift-image-registry.svc:5000/
openshift/bitnami-nginx:1.26 --replicas=3 --dry-run=client -o yaml > nginx-
deploy.yml
oc apply -f nginx-deploy.yml
```

Let's create the ClusterIP Internal service declaratively

```
oc expose deploy/nginx --type=ClusterIP --port=8080 --dry-run=client -o yaml
oc expose deploy/nginx --type=ClusterIP --port=8080 --dry-run=client -o yaml >
nginx-clusterip-svc.yml
oc apply -f nginx-clusterip-svc.yml
oc get svc
oc describe svc/nginx
```

Let's test the clusterip internal service using the test pod we created

```
oc rsh pod/test-pod
curl http://nginx:8080 # Service discovery - accessing service by its name
curl http://nginx.jegan-project.svc.cluster.local:8080 # Recommended compared to
previous command
curl http://172.30.240.164:8080
```

Lab - Creating an external NodePort service in declarative style

```
oc project jegan-project

# Delete existing clusterip service before creating nodeport service
oc delete -f nginx-clusterip-svc.yml

oc expose deploy/nginx --type=NodePort --port=8080 --dry-run=client -o yaml
oc expose deploy/nginx --type=NodePort --port=8080 --dry-run=client -o yaml >
nginx-nodeport-svc.yml
oc apply -f nginx-nodeport-svc.yml
oc get svc
oc describe svc/nginx # Get the NodePort from this command and substitute below

oc get nodes -o wide
```

Let's test the nodeport external service (don't need to run this inside pod shell)

```
curl http://192.168.100.11:31728 # Master1 IP
curl http://192.168.100.12:31728 # Master2 IP
curl http://192.168.100.13:31728 # Master3 IP
curl http://192.168.100.21:31728 # Worker1 IP
curl http://192.168.100.22:31728 # Worker2 IP
curl http://192.168.100.23:31728 # Worker3 IP

curl http://master01.ocp4.palmeto.org:31728
curl http://master02.ocp4.palmeto.org:31728
curl http://master03.ocp4.palmeto.org:31728
curl http://worker01.ocp4.palmeto.org:31728
curl http://worker02.ocp4.palmeto.org:31728
curl http://worker03.ocp4.palmeto.org:31728
```

Lab - LoadBalancer service in declarative style

```
oc project jegan-project

# Delete existing node port service
oc delete -f nginx-nodeport-svc.yml

# Create the LoadBalancer service in declarative style
oc expose deploy/nginx --type=LoadBalancer --port=8080 --dry-run=client -o yaml >
nginx-lb-svc.yml
oc apply -f nginx-lb-svc.yml

oc get svc
oc describe svc/nginx # You need to find your loadbalancer service ip and use it
below in the curl

# Testing the loadbalancer service
curl http://192.168.100.50:8080
```

Lab - Rolling update (upgrading nginx from version 1.26 to 1.27 and later to 1.28)

Let's delete your existing project

```
oc delete project jegan-project
```

Let's create new project

```
oc new-project jegan-project
oc create deploy nginx --image=image-registry.openshift-image-registry.svc:5000/
openshift/bitnami-nginx:1.26 --replicas=3 --dry-run=client -o yaml > nginx-
deploy.yml
oc apply -f nginx-deploy.yml

oc get pods -o yaml | grep image
```

Lab - Create an external route for nginx service

```
# Delete existing project
oc delete project jegan-project

# Deploy nginx
oc new-project jegan-project
oc create deploy nginx --image=image-registry.openshift-image-registry.svc:5000/
openshift/bitnami-nginx:1.26 --replicas=3 --dry-run=client -o yaml > nginx-
deploy.yml
oc apply -f nginx-deploy.yml

# Create clusterip internal service for nginx deployment
oc expose deploy/nginx --port=8080 --dry-run=client -o yaml > nginx-internal-
svc.yml
oc apply -f nginx-internal-svc.yml

# Create an external route for internal service
oc expose svc/nginx
oc get route
curl http://nginx-jegan-project.apps.ocp4.palmeto.org
```

Info - ConfigMap

- is a map that stores data in key/value format
- instead of hard-coding, we store certain details like NFS Server IP, NFS Share Path, etc in the config map
- generally used for storing this like
 - JDK_HOME, MAVEN_HOME, LOG_PATH, etc.,
- deployment can retrieve the details from the ConfigMap and use it
- this is good for non-sensitive data

Info - Secrets

- is also map that stores data in key/value format
- it is used to store any sensitive confidential data like login credentials, certificates, etc.,

Info - Persistent Volume (PV)

- is an external disk/storage
- this type of storage can come from NFS,NAS,AWS S3, etc.,
- this is always created on the cluster-wide and accessible to all projects in the cluster
- this can be provisioned manually by Administrators or can be provisioned on demand dynamically if there is a StorageClass
- Generally in case of NFS
 - we need to mention NFS Server IP/Hostname
 - Disk size required in MiB/GiB/TiB
 - StorageClass (Optional)
 - Access
 - ReadWriteOnce
 - ReadWriteMany

Info - Persistent Volume Claim (PVC)

- a application that runs in a Pod will have to request for external storage in Openshift/Kubernetes by defining its requirement in a Persistent Volume Claim
- generally created by non-admin, usually the dev team and it is created in the project scope
- If the Storage Controller is able to locate a matching Persistent Volume, then it let the PVC go and claim and use the PV
- the applicaiton Pod that needs the store will refer the PVC name in the deployment, and it can request the external store to be mounted in its preferred mount point within the Pod

Lab - Deploy wordpress and mariadb multi-pod application

```
# Clone the TekTutor Training repository if you haven't done it already on the lab machine
```

```

cd ~
git clone https://github.com/tektutor/openshift-july-2026.git
cd openshift-july-2026
cd Day3/wordpress-with-configmaps-and-secrets
# Ensure name 'jegan' is replaced with yours in all the yml file
# Ensure /var/nfs/jegan/wordpress, /var/nfs/jegan/mysql is replaced with your
linux user ( update mysql-pv.yml mysql-pvc.yml wordpress-pv.yml wordpress-
pvc.yml )
# Ensure NFS IP is updated to 192.168.10.201 in case you are working in server 2
# Ensure mysql-deploy.yml update the pvc name with the one that you created
# Ensure wordpress-deploy.yml update the pvc name with the one that you created

oc delete project jegan-project
oc new-project jegan-project
./deploy.sh

oc get pods

# Then switch to Openshift webconsole Topolgy and click on the wordpress route(up
arrow) to see the blog

```

Demo - Installing Helm package manager

```

sudo su -
curl -fsSL -o get_helm.sh https://raw.githubusercontent.com/helm/helm/main/
scripts/get-helm-4
chmod 700 get_helm.sh
./get_helm.sh

```

Info - Helm Overview

- helm is a package manager for Kubernetes & Openshift
- helm is a Kubernetes/Openshift aware tool, hence it knows how to find the dependency and figure out the sequence in which K8s/Openshift resources must be created
- the helm packaged application is called Helm chart
- Openshift comes with pre-integrated Helm support
- However, we can deploy application from Helm Charts from command-line as well as from the Openshift webconsole

Lab - Packaging wordpress multi-pod application as Helm chart and deploying into Openshift

```
cd ~
mkdir -p wordpress-helm-chart
cd wordpress-helm-chart

helm version
helm create wordpress
tree
cd wordpress/templates
rm *
rm -rf tests
cp ~/openshift-july-2026/Day3/wordpress-helm-chart/manifest-scripts/*.yaml .
cd ..
cp ~/openshift-july-2026/Day3/wordpress-helm-chart/values.yaml .
cd ..
tree wordpress

helm package wordpress
helm install worpress wordpress-0.1.0.tgz
oc get pods

# Switch to your openshift webconsole topology in your project

# Once you are done you may undeploy
helm list
helm uninstall wordpress
```

Day 4

Demo - Statefulset - mysql cluster

```
cd ~/openshift-july-2026
git pull
cd Day4/statefult
oc apply -f sfs.yml
oc get pods -w
```

Lab - Node Affinity

Note

- There are two types of Node Affinity
 1. Preferred
 - In this case, the scheduler will try to locate nodes that meets the criteria mentioned under preferred node affinity section, if scheduler finds such nodes, your application gets deployed there
 - In case the scheduler is not able to find nodes that meets your criteria, the pods gets deployed in any node
 2. Required
 - In this case, the scheduler will only deploy your pods on nodes that meets the criteria
 - Until scheduler finds such nodes that meets your criteria, your pods will be in the Pending state

Let's try the preferred node affinity

```
cd ~/openshift-july-2026
git pull
cd Day4/node-affinity
cat preferred-node-affinity.yml
oc project jegan-project

# List all nodes that has label disk=ssd
oc get nodes -l disk=ssd

# Assuming node nodes meets that criteria, let's see what scheduler does in this
scenario
oc apply -f preferred-node-affinity.yml
oc get pods -o wide # As the scheduler couldn't find nodes that has label
disk=ssd, it would have deployed your pods as usual

# Let's label a node with lab disk=ssd
oc label node/worker03.ocp4.palmeto.org disk=ssd
oc get nodes -l disk=ssd

# List and see to note the pods deployed already will not be redeployed to
worker03 as node affinity is not applied on already
# deployed pods
oc get pods -o wide

# Let's delete the deployment and redeploy
oc delete -f preferred-node-affinity.yml
oc apply -f preferred-node-affinity.yml
```



```

oc get pods -o wide # Now all your nginx pods must be deployed to worker03

# Let's delete the preferred-node-affinity.yml deployment to proceed with
required-node-affinity
oc delete -f preferred-node-affinity.yml

# Let's remove the label from worker03 node
oc label node/worker03.ocp4.palmeto.org disk-
oc get nodes -l disk=ssd

# Let's deploy the required-node-affinity nginx deployment
oc apply -f required-node-affinity.yml
oc get pods -o wide # all pods are supposed to be in Pending as scheduler
couldn't find a node that meets your app required criteria

# Label one of the worker node with disk=ssd
oc label node/worker01.ocp4.palmeto.org disk=ssd
oc get pods -o wide # now all your nginx pods should be running in worker01

```

Lab - Ingress

Note

- Ingress is a way to make our application accessible outside the cluster with public url
- Ingress is not Service, it is routing(forwarding) set of rules
- In your openshift cluster, depending on which Load Balancer is setup by your Openshift Cluster Administrator, we need to use either
 1. HAProxy Ingress Controller or
 2. Nginx Ingress Controller or
 3. F5 Ingress Controller, etc
- Whenever we create an Ingress resource, the Ingress Controller that is active in the cluster will keep monitoring for the Ingress resource created in any project within the cluster
- Whenever the Controller detects such a Route, it picks the rules mentioned in the Ingress resource and configures a Load Balancer
- HAProxy Ingress Controller only know how to configure a HAProxy Load Balancer
- Nginx Ingress Controller knows only how to configure a Nginx Load Balancer
- It is also important that our Ingress annotation mentions the Ingress Controller that should manage our Ingress
- Ingress is not a service
 - as Ingress will forward the call to one of the service based on rule defined in the Ingress

- in other words, the backend of Ingress is set of services (ClusterIP, NodePort or LoadBalancer Services)
- Ingress represents a group of Services
- Service (ClusterIP, NodePort or LoadBalancer)
- Service represents a group of load-balanced pods
- the backend of these services are Pods

Let's delete and recreate a new project

```
oc delete project jegan-project
oc new-project jegan-project

cd ~/openshift-july-2026
git pull
cd Day4/ingress
oc apply -f nginx-deploy.yml
oc apply -f hello-deploy.yml
oc apply -f nginx-svc.yml
oc apply -f hello-svc.yml
oc describe svc/nginx
oc describe svc/hello
oc apply -f ingress.yml
oc get ingress

curl http://tektutor.apps.ocp4.palmeto.org/hello
curl http://tektutor.apps.ocp4.palmeto.org/nginx
```

Lab - How two microservices exchange data via Queue in a asynchronous fashion

```
oc delete project jegan
oc new-project jegan

oc new-app --name=jms-producer https://github.com/tektutor/openshift-july-2026.git --context-dir=Day4/jms-demo/producer --strategy=docker
oc new-app --name=jms-consumer https://github.com/tektutor/openshift-july-2026.git --context-dir=Day4/jms-demo/consumer --strategy=docker

oc logs -f buildconfig/jms-producer
oc logs -f buildconfig/jms-consumer

oc get pods

# Terminal 1
oc logs -f deploy/jms-producer
```

```
# Terminal 2
oc logs -f deploy/jms-consumer
```

Lab - Securing your application using https url (edge route)

```
oc delete project jegan-project
oc new-project jegan-project
oc create deploy nginx --image=image-registry.openshift-image-registry.svc:5000/
openshift/bitnami-nginx:1.28 --replicas=3
oc get deploy,pods

# The command below creates a clusterip internal service for nginx deployment
oc expose deploy/nginx --port=8080

openssl version
# Let's generate a private key
openssl genrsa -out key.key

# Let's create a public key using the private key for your organization domain
openssl req -new -key key.key -out csr.csr -subj="/CN=nginx-
jegan.apps.ocp4.palmeto.org"

# Sign the public key using the private key and generate certificate(.crt)
openssl x509 -req -in csr.csr -signkey key.key -out crt.crt

# The command below creates an edge(https) route
oc create route edge --service nginx --hostname nginx-jegan.apps.ocp4.palmeto.org
--key key.key --cert crt.crt

oc get route
curl --insecure https://nginx-jegan.apps.ocp4.palmeto.org
```

Day 5

Info - How many users one Pod can handle?

- there is no fixed number
- it depends on your application's per-request cost and its concurrency model
- how to compute the concurrent capacity
$$\text{concurrent requests} = (\text{throughput in req/sec}) \times (\text{average request duration in sec})$$

- if each request takes 200ms and one pod can support upto 250~500 request/second
- a lightweight SpringBoot application with no database dependency with 500m CPU and 512Mi RAM can approximately handle 1000+ requests/second
- the same springboot pod with 40ms database query per request drops to roughly 200~400 concurrent users
- Tomcat can handle 10 concurrent database-bound requests
- Benchmarking is the right way to find the number as concurrency depends on your application's ability and how long it takes to respond a single request

Info - Openshift S2I

- In Kubernetes, we can only deploy application using container images
- In Openshift, we can deploy applicaiton using container images and from github(or any other version control) url i.e from source code
- This feature in Openshift is called S2I (Source to Image)
- Openshift supports many different types of strategies
 1. Docker
 2. Source
 3. Custom
 4. Pipeline(Jenkins/TekTon)
 5. Binary(S2I Binary)

Info - Openshift S2I Docker Strategy

- You will have to provide your version control ur, for example - <https://github.com/tektutor/spring-ms.git>
- The GitHub url should have your application source code and Dockerfile
- BuildConfig will be autogenerated by Openshift
- Application will be build to create application binary, using your application binary it will create a custom image
- the custom image built will be pushed into Openshift's Internal Image Registry
- It will then deploy your application using the image it pushed into Openshift Internal Image Registry
- It will automatically create BuildConfig, ImageStream, Deployment and Service for your application
- `oc expose svc/hello` - you need to create your own route
- you can then access the application using the route url from command-line or from your openshift web console

Info - OpenShift S2I Source Strategy

- You will have to provide your version control url, for example - `https://github.com/tektutor/spring-ms.git`
- The GitHub url should have your application source code, Dockerfile is not required even if it is there it will be ignored
- BuildConfig will be autogenerated by Openshift
- Application will be build to create application binary, using your application binary it will create a custom image
- the custom image built will be pushed into Openshift's Internal Image Registry
- It will then deploy your application using the image it pushed into Openshift Internal Image Registry
- It will automatically create Dockerfile, BuildConfig, ImageStream, Deployment and Service for your application
- `oc expose svc/hello` - you need to create your own route
- you can then access the application using the route url from command-line or from your openshift web console

Info - Microservice Overview

- Microservice is an architectural style where application is built as a collection of small, independent and loosely coupled services
- each microservice will be responsible for specific business logic
- microservice to microservice communications can be achieved with HTTP/REST, gRPC or message queues
- Key properties of Microservices
 - Single Responsibility
 - Can be deployed independently
 - Technology agnostic - can be developed using any language stack
 - Decentralized Database
 - Resilience - Failures in one service will not affect others or will not bring down the entire application
 - Scalability - every microservice can be scaled up/down independent of other microservices

Info - Monolithic Application Overview

- is a traditional application that is built as a single application binary
- all components will be bundled in the single application binary
 - Frontend

- Business Layer
- Data Access Layer
- Key properties of Monolithic applications
 - Single codebase
 - Tightly coupled components
 - Entire application is deployed as a single unit
 - Centralized Database
 - Scaling Limitations

Info - Secure your Openshift Cluster and applications deployed in Openshift

Authentication & Authorization

- Authentication
 - Could setup Identity providers (IDPs) to authenticate openshift login
 - LDAP, GitHub, GitLab, Google, HTPasswd, Keycloak (OIDC)
 - ServiceAccounts vs User accounts
 - Token and OAuth flows (oc whoami -t)
- Authorization
 - RBAC (Role-Based Access Control)
 - Roles, ClusterRoles, RoleBindings, ClusterRoleBindings
 - Custom Roles vs. default roles (admin, view, edit, etc.)
 - Least privilege principle

Network Security

- Pod-to-Pod & Pod-to-Service
- NetworkPolicy (namespaced) and ClusterNetworkPolicy (cluster-wide)
- Default deny-all and zero-trust networking
- Controlling ingress/egress with egress firewall rules
- Ingress Security
- Routes and TLS termination types
- Securing routes with certificates

Cluster Security & Hardening

- Security Context Constraints (SCC) – OpenShift equivalent of PodSecurityPolicies

- restricted, anyuid, privileged SCCs
- Creating custom SCCs
- Admission Controllers and Validating/Mutating Webhooks
- Pod Security Standards (Baseline, Restricted)
- Running containers as non-root
- Disabling privileged escalation

Image & Supply Chain Security

- ImageStreams and trusted registries
- image signature verification (oc adm policy add-scc-to-user, oc adm signature)
- Image scanning (e.g., Red Hat Quay, Clair, Trivy)
- Content Trust and ImagePolicy admission
- Immutable images and reproducible builds
- Security scanning in CI/CD pipelines

Secrets & Config Security

- Kubernetes Secrets vs ConfigMaps
- Encryption at rest with KMS providers (e.g., AWS KMS, Vault)
- Sealed Secrets (Bitnami SealedSecrets)
- External secret managers (HashiCorp Vault, CyberArk, AWS Secrets Manager)
- Securing environment variables in deployments

Platform Access & API Security

- API server authentication and RBAC for oc and kubectl users
- OpenShift Console access restrictions
- API audit logging
- Certificate management for API server and etcd
- OAuth token lifetimes and rotation

Compliance & Monitoring

- Compliance Operator (CIS Benchmark for OpenShift)
- OpenSCAP scans and remediation
- Audit logging and forwarding (e.g., to ELK or SIEM)

- Security events in oc get events and cluster logs
- Cluster-wide alerts for security breaches

Security Tools & Operators

- Compliance Operator – CIS, NIST profiles
- Gatekeeper / OPA – policy as code
- ACS (Red Hat Advanced Cluster Security) – runtime threat detection, network graph
- Falco – behavioral threat detection
- Trivy – image scanning in CI/CD

Secure Multitenancy

- Namespace isolation best practices
- Project quotas and limits
- Network segmentation per tenant
- RBAC delegation to team leads
- Resource constraints to prevent noisy neighbor issues

Security Updates & Patch Management

- OpenShift upgrade process & channels (fast, stable)
- Node OS updates (RHCOS)
- Operator lifecycle management (OLM) security
- CVE tracking and remediation using Red Hat Insights

Runtime & Workload Security

- Pod security context:
 - runAsUser, fsGroup, readOnlyRootFilesystem
- Dropping Linux capabilities
 - Seccomp profiles and AppArmor (on supported platforms)
- Container resource limits to prevent DoS
- Detecting crypto-mining and lateral movement

Ingress/Egress & Perimeter Security

- Securing ingress with TLS, WAF integration (e.g., F5, HAProxy)
- Rate limiting & DoS protection
- Egress control with EgressFirewall and proxies
- Bastion hosts & jump boxes for controlled access

Backup & Disaster Recovery (Security Angle)

- etcd encryption and backup security
- Secrets encryption during backups
- Immutable backup storage
- RBAC restrictions for restore operations

Zero Trust Architecture

- Network segmentation
- Identity-aware access to workloads
- Mutual TLS between services
- Policy enforcement using OPA/Gatekeeper
- Just-in-time access control

Best Practices & Standards

- CIS Benchmark for OpenShift
- NIST 800-53 alignment
- ISO 27001 security practices
- Regular pen-testing and red teaming
- Shift-left security in CI/CD

Info - Jenkins Overview

- Kohsuke Kawaguchi was a former employee of Sun Microsystems
- In he developed a Continuous Integration Build Server to automate his testing at Sun Microsystems as a hobby
- This product originally called Hudson and was an opensource project
- Meanwhile, Oracle acquired Sun Microsystems

- Koshuke Kawaguchi founded a company called Cloudbees
- When he quit Oracle, he had created a fork of Hudson branch and new branch was named as Jenkins
- Cloudbees is the Enterprise variant of Jenkins
- Jenkins has over 10000+ active opensource contributors around the globe

Info - Openshift Network Model

- Openshift's network operates in layers
- It is important to understand how traffic flows through each layers for debugging & troubleshooting
- Core Concepts
 1. Container Network Interface (CNI)
 2. Pod to Pod Communication
 3. Services and Load Balancing

Container Network Interace (CNI)

- Openshift uses a CNI plugin to handle pod networking
- the default is OVN-Kubernetes (Openshift Virtual Network)
- OVN creates a virtual network overlay where every pod gets a real, routable IP address
 - from the cluster subnet

Pod to Pod Communication

- When a Pod A sends traffic to Pod B, OVN handles the routing without needing kube-proxy
- Packets flow through the overlay network
- Both pods live on the same logical network namespaces, if even if they are on different nodes
- The network plugin encapsulates packets and routes them to the correct node

Services and Load Balancing

- Services abstract pods IP addresses as Pod IPs are ephemeral
- When you create a service, it gets a stable cluster IP
- The kube-proxy watches endpoints and programs iptable rules on every node
- The rules use DNAT with random probability-weighted distribution for

new connections

- each new connection picks a backend pod randomly according to their weights
- Pod to Pod communication via Internal Service
 - Pod A queries the service DNS name, gets resolved to cluster IP
 - Packet hits a node's iptable rules
 - DNAT rewrites the destination to a real Pod IP
 - OVN routes to the target node and pod
- Pod to Pod communication via External Route
 - Uses Route or Ingress object
 - Route creates a HAProxy rule on the ingress controller pods
 - Traffic lands on the Ingress controller, which forwards to the service
 - service load balancing takes over from there
- Egress (pod to external traffic)
 - Pods use their Node's IP as source
 - Outbound traffic appears to come from the cluster, not from individual pods

Info - Subnet

- Range of IP address
- a bigger network is broken down to small network called subnet to apply different security policy or access
- IPV4 IP Address
 - 100.200.3.20 (4 bytes)
 - A.B.C.D
 - A - 1 byte (8 bits)
 - B - 1 byte (8 bits)
 - C - 1 byte (8 bits)
 - D - 1 byte (8 bits)
 - 1 byte represents numbers in the range 0 to 255 (256 values)
- Let's say there is subnet 10.131.0.0/24
 - 10 is fixed (8 bits)
 - 131 is fixed (8 bits)
 - 0 is fixed (8 bits)
 - How many IP addresses are there in 10.131.0.0/24 - 256 IP addresses are there

Flannel Network Model

Flannel is a simple and easy way to configure a layer 3 network fabric designed for Kubernetes.

It provides networking for container clusters by creating a virtual network that spans across all nodes.

Key Concepts

- Overlay Network
- Creates a virtual network layer on top of the existing host network
- Each node gets a subnet from a larger cluster network
- Pods can communicate across nodes using this overlay
- Network Backends
- VXLAN (Default)

Encapsulates packets in UDP

- Works across most network infrastructures
- Good performance with modern kernels


Host-Gateway

- Uses host routing tables
- Better performance but requires Layer 2 connectivity
- No packet encapsulation overhead

UDP

- Legacy backend
- User-space packet forwarding
- Lower performance but maximum compatibility

Architecture



[illegible]

```
| Routing Table (Node A):
B): |
```

| Routing Table (Node

Src MAC: 02:42:0a:f4:01:0a (Pod1 eth0)
EtherType: 0x0800 (IPv4)
IP Header
Version: 4, Header Length: 20 bytes
Source IP: 10.244.1.10
Destination IP: 10.244.2.10
Protocol: 1 (ICMP)
TTL: 64
ICMP Header
Type: 8 (Echo Request)
Code: 0
Identifier: 1234
Sequence: 1
Data: "Hello from Pod1"

Packet on Node A (before VXLAN encapsulation)

```
# On Node A - Capture on cni0 bridge
sudo tcpdump -i cni0 -nn icmp
# 15:30:45.123456 IP 10.244.1.10 > 10.244.2.10: ICMP echo request, id 1234, seq 1
```

VXLAN Encapsulated Packet on Physical Network

```
# On Node A - Capture on eth0 (physical interface)
sudo tcpdump -i eth0 -nn 'port 8472'
```

Encapsulated Packet Structure:

Outer Ethernet Header
Dst MAC: aa:bb:cc:dd:ee:02 (Node B eth0)
Src MAC: aa:bb:cc:dd:ee:01 (Node A eth0)
EtherType: 0x0800 (IPv4)
Outer IP Header
Version: 4, Header Length: 20 bytes
Source IP: 10.0.1.100 (Node A)
Destination IP: 10.0.1.101 (Node B)
Protocol: 17 (UDP)

TTL: 64
UDP Header
Source Port: 45678 (random) Destination Port: 8472 (VXLAN) Length: 78 bytes
VXLAN Header
Flags: 0x08 (Valid VNI) VNI: 1 (Virtual Network Identifier) Reserved: 0x000000
Inner Ethernet Header
Dst MAC: 02:42:0a:f4:02:01 (Node B cni0) Src MAC: 02:42:0a:f4:01:01 (Node A cni0) EtherType: 0x0800 (IPv4)
Inner IP Header
Version: 4, Header Length: 20 bytes Source IP: 10.244.1.10 (Pod1) Destination IP: 10.244.2.10 (Pod3) Protocol: 1 (ICMP) TTL: 63 (decremented by 1)
ICMP Header
Type: 8 (Echo Request) Code: 0 Identifier: 1234 Sequence: 1 Data: "Hello from Pod1"

Packet Capture Commands and Output

On Node A:

<pre># Capture original packet leaving Pod1 sudo tcpdump -i veth12345678 -nn icmp -v # 15:30:45.123456 IP (tos 0x0, ttl 64, len 84) 10.244.1.10 > 10.244.2.10: # ICMP echo request, id 1234, seq 1, length 64 # Capture VXLAN encapsulated packet on physical interface sudo tcpdump -i eth0 -nn 'udp port 8472' -v # 15:30:45.123460 IP (tos 0x0, ttl 64, len 134) 10.0.1.100.45678 > 10.0.1.101.8472:</pre>
--


```
# VXLAN, flags [I] (0x08), vni 1
# IP (tos 0x0, ttl 63, len 84) 10.244.1.10 > 10.244.2.10:
# ICMP echo request, id 1234, seq 1, length 64
```

On Node B:

```
# Capture incoming VXLAN packet
sudo tcpdump -i eth0 -nn 'udp port 8472' -v
# 15:30:45.123465 IP (tos 0x0, ttl 64, len 134) 10.0.1.100.45678 >
10.0.1.101.8472:
# VXLAN, flags [I] (0x08), vni 1
# IP (tos 0x0, ttl 63, len 84) 10.244.1.10 > 10.244.2.10:
# ICMP echo request, id 1234, seq 1, length 64

# Capture decapsulated packet on cni0
sudo tcpdump -i cni0 -nn icmp -v
# 15:30:45.123470 IP (tos 0x0, ttl 63, len 84) 10.244.1.10 > 10.244.2.10:
# ICMP echo request, id 1234, seq 1, length 64
```

Indide Pod3

```
# Final packet received by Pod3
tcpdump -i eth0 -nn icmp -v
# 15:30:45.123475 IP (tos 0x0, ttl 63, len 84) 10.244.1.10 > 10.244.2.10:
# ICMP echo request, id 1234, seq 1, length 64
```

Return Path (Pod3 → Pod1)

The ICMP reply follows the reverse path:

```
# Pod3 sends reply
# 15:30:45.123480 IP 10.244.2.10 > 10.244.1.10: ICMP echo reply, id 1234, seq 1

# Node B encapsulates and sends via VXLAN
# 15:30:45.123485 IP 10.0.1.101.54321 > 10.0.1.100.8472:
# VXLAN, flags [I] (0x08), vni 1
# IP 10.244.2.10 > 10.244.1.10: ICMP echo reply, id 1234, seq 1

# Pod1 receives the reply
# 15:30:45.123490 IP 10.244.2.10 > 10.244.1.10: ICMP echo reply, id 1234, seq 1
```

Info - Calico Openshift Network Plugin

- Calico Network in OpenShift
- Calico is a popular Container Network Interface (CNI) plugin that

provides networking and network security for containerized applications in OpenShift clusters.
What is Calico?

- Calico is an open-source networking solution that delivers:

Pod-to-pod networking across nodes

- Network policy enforcement for security
- IP address management (IPAM)
- Service mesh integration
- Key Features in OpenShift
 1. Pure Layer 3 Networking
 - Uses standard IP routing instead of overlay networks
 - Better performance with lower latency
 - Simplified troubleshooting using standard networking tools
 2. Network Policies
 - Kubernetes-native network policies
 - Advanced Calico network policies with additional features
- Ingress and egress traffic control
 - Application-layer policy enforcement
- Scalability
 - Supports thousands of nodes
 - Efficient BGP routing protocols
 - Distributed architecture without single points of failure
- Security
 - Workload isolation at the network level
 - Encryption in transit
 - Integration with service mesh security
- Architecture Components
 - Felix Agent
 - Runs on each node as a DaemonSet
 - Programs routing tables and iptables rules
 - Handles network policy enforcement
 - BGP Client (BIRD)
 - Distributes routing information between nodes
 - Maintains network topology
 - Enables cross-node pod communication
 - Calico Controller
 - Watches Kubernetes API for network policy changes
 - Translates policies into Felix-readable format
 - Manages IP address allocation
- Common Use Cases
 - Multi-tenant environments requiring strong network isolation
 - High-performance applications needing low network latency
 - Hybrid cloud deployments with consistent networking policies
 - Compliance requirements demanding network traffic control

Comparison

Feature	Calico	OpenShift SDN	OVN-Kubernetes
Performance	High	Medium	Medium-High
Network Policies	Advanced	Basic	Basic
Complexity	Medium	Low	Medium
BGP Support	Yes	No	No
Overlay Network	Optional	Yes	Yes
eBPF Support	Yes	No	Limited
IPv6 Support	Yes	Limited	Yes
Windows Support	Yes	No	Yes
Encryption	WireGuard	IPSec	IPSec
IPAM	Custom	Built-in	Built-in
Troubleshooting	Standard	OpenShift	OVN Tools
	Tools	Tools	

Info - OpenShift Weave Network Plugin

- Weave Network Plugin in OpenShift
 - Weave Network (Weave Net) is a Container Network Interface (CNI) plugin that
 - can be used with OpenShift to provide networking between containers and pods across a cluster.
- What is Weave Network?
 - Weave Network creates a virtual network that connects Docker containers across multiple hosts
 - and enables their automatic discovery. It provides:
- Overlay Network: Creates a Layer 2 network overlay that spans multiple hosts
- Automatic IP Management: Assigns IP addresses to containers automatically
- Service Discovery: Built-in DNS-based service discovery
- Network Segmentation: Support for network policies and microsegmentation
- Key Features in OpenShift Context
 1. Multi-Host Networking
 - Connects pods across different OpenShift nodes
 - Handles pod-to-pod communication seamlessly
 - Supports both IPv4 and IPv6
 2. Network Policies
 - Implements Kubernetes NetworkPolicy resources
 - Provides microsegmentation capabilities
 - Allows traffic filtering between pods/namespaces
 3. Encryption
 - Optional encryption of inter-host traffic

- Uses NaCl crypto library for performance
 - Helps secure pod communication across untrusted networks
4. Integration with OpenShift
- Works as a CNI plugin
 - Integrates with OpenShift's Software Defined Networking (SDN)
 - Compatible with OpenShift service mesh

Architecture Components

- Weave Router
 - Runs as a DaemonSet on each node
 - Handles packet forwarding and routing
 - Maintains network topology information
- Weave Net Plugin
 - CNI plugin that interfaces with container runtime
 - Allocates IP addresses to pods
 - Configures network interfaces
 - IPAM (IP Address Management)
 - Distributes IP address ranges across nodes
 - Prevents IP conflicts
 - Supports custom IP ranges

Useases

- Multi-cloud deployments where encryption is required
- Strict network segmentation requirements
- Legacy applications that need specific networking features
- Development environments requiring flexible networking
- Comparison with OpenShift SDN

Limitations

- Performance overhead especially with encryption enabled
- Additional complexity compared to default OpenShift SDN
- Resource consumption from running additional networking components
- Troubleshooting complexity due to overlay networking
- Weave Network provides a robust networking solution for OpenShift, particularly useful when advanced networking features like - encryption and comprehensive network policies are required.

Flannel vs Calico vs Weave

+-----+-----	+-----+-----
+-----+-----	+-----+-----
Feature	Calico

Flannel	Weave	
-----+	-----+	-----+
Overlay Technology	BGP (L3), VXLAN, IP	VXLAN, UDP, host-
gw VXLAN, UDP		
-----+	-----+	-----+
Network Policies	Full NetworkPolicy +	No native
support Full NetworkPolicy		
	Calico policies	
	support	
-----+	-----+	-----+
Encryption	WireGuard, IPSec	No native
encryption Built-in NaCl encryption		
-----+	-----+	-----+
Performance	Excellent (native BGP)	Good (simple
overlay) Good (with encryption		
	overhead)	
-----+	-----+	-----+
Setup Complexity	Complex (BGP knowledge	
Simple Moderate		
	required)	
-----+	-----+	-----+
IP Address Management	IPAM with IP pools	Simple host-
subnet Distributed IPAM		
allocation		
-----+	-----+	-----+
Service Discovery	Standard Kubernetes DNS	Standard
Kubernetes DNS Built-in DNS + K8s DNS		
-----+	-----+	-----+
Multi-tenancy	Advanced policy engine	Basic
namespace Network segmentation		
isolation		
-----+	-----+	-----+
Scalability	Excellent (BGP routing)	Good (limited by
etcd) Good (mesh topology)		
-----+	-----+	-----+
Cross-subnet Support	Native (BGP)	Limited

(requires mode)	Good (automatic discovery)	host-gw
+-----+-----+		
Windows Support Limited	Yes No	
+-----+-----+		
eBPF Support No	Yes (advanced features) No	
+-----+-----+		
Resource Usage low	Low to moderate Moderate to high	Very
+-----+-----+		
Troubleshooting Simple	Complex (BGP knowledge) Moderate	
+-----+-----+		
Enterprise Features Limited	Calico Enterprise Weave Cloud (deprecated) (advanced security)	
+-----+-----+		
IPv6 Support Basic	Full dual-stack Yes	
+-----+-----+		
Network Observability metrics	Flow logs, metrics Network map, metrics	Basic
+-----+-----+		
Maturity Mature	Very mature Mature	
+-----+-----+		
Best Use Cases deployments	- Large scale clusters - Multi-cloud setups - Advanced security - Encryption required - Compliance requirements - Easy troubleshooting - Multi-cloud - Legacy app support	- Simple - Learning/ - Resource
+-----+-----+		
Vendor/Maintainer	Tigera (formerly	CoreOS/Red

Hat	Weaveworks	
	Project Calico)	
+-----+-----+		
+-----+-----+		
OpenShift Support	Yes (OVN-Kubernetes	
Limited	Third-party plugin	
	uses Calico for policies)	
+-----+-----+		
+-----+-----+		

Lab - Deploying our custom spring-boot app into Openshift using S2I docker strategy

```
oc new-app --name=hello-microservice https://github.com/tektutor/openshift-july-2026.git --context-dir=Day5/simple-springboot-microservice --strategy=docker

oc expose svc/hello-microservice

oc logs -f bc/hello-microservice

oc get svc,route

curl --insecure http://hello-microservice-jegan.apps.ocp4.palmeto.org
```

Lab - Horizontal Pod Auto-scaling based on CPU Utilization

```
oc delete project jegan
oc new-project jegan

cd ~
git clone https://github.com/tektutor/openshift-july-2026.git
cd openshift-july-2026
cd Day5/auto-scaling
oc create -f hello-deploy.yml --save-config=true
oc get pods
oc create -f hello-hpa.yml --save-config

oc expose deploy/nginx --port=8080
oc expose svc/nginx
oc get route
```

We need to stress the pod with more traffic

```
ab -k -n 200000 -c 1000 https://nginx-jegan.apps.ocp4.palmeto.org/
```

Lab - CI/CD with Jenkins & Red Hat OpenShift

The screenshot shows the Red Hat OpenShift console interface. The left sidebar contains navigation links for Home, Overview, Projects, Search, API Explorer, Events, Favorites, Ecosystem, Software Catalog, Installed Operators, Helm, Workloads, Networking, Services, Routes, and Ingresses. The main content area is titled 'Overview' and displays cluster details. A notification at the top states: 'You are logged in as a temporary administrative user. Update the cluster OAuth configuration to allow others to log in.'

Getting started resources

- Set up your cluster**: Finish setting up your cluster with recommended configurations. Links include 'Take console tour', 'Add identity providers', 'Configure alert receivers', and 'View all steps in documentation'.
- Build with guided documentation**: Follow guided documentation to build applications and familiarize yourself with key features. Links include 'Enable the Developer Perspective' and 'Impersonating the system:admin user'.
- Explore new features and capabilities**: Links to 'OpenShift AI', 'Trusted Software Supply Chain', and 'OpenShift Lightspeed'.

Details (View settings)

- Cluster API address**: <https://api.ocp4.palmeto.org:6443>
- Cluster ID**: 07cbd7d8-2763-497c-8763-cccc67e8be8
- [OpenShift Cluster Manager](#)

Status (View alerts)

- Cluster**: Active
- Control Plane**: Active
- Operators**: Active
- Dynamic Plugins**: Active
- Insights**: 0 issues found
- TargetDown**: Jun 14, 2026, 8:39 PM (View details)

Activity

- Ongoing**: There are no ongoing activities.
- Recent events**: 2:16 PM Failed to apply de... 7:16 PM Failed to apply de...



The screenshot shows the Red Hat OpenShift console interface, specifically the 'Projects' page. The left sidebar is the same as the previous screenshot. The main content area is titled 'Projects' and displays a list of projects. A notification at the top states: 'You are logged in as a temporary administrative user. Update the cluster OAuth configuration to allow others to log in.'



Projects (Create Project)



Filter by name


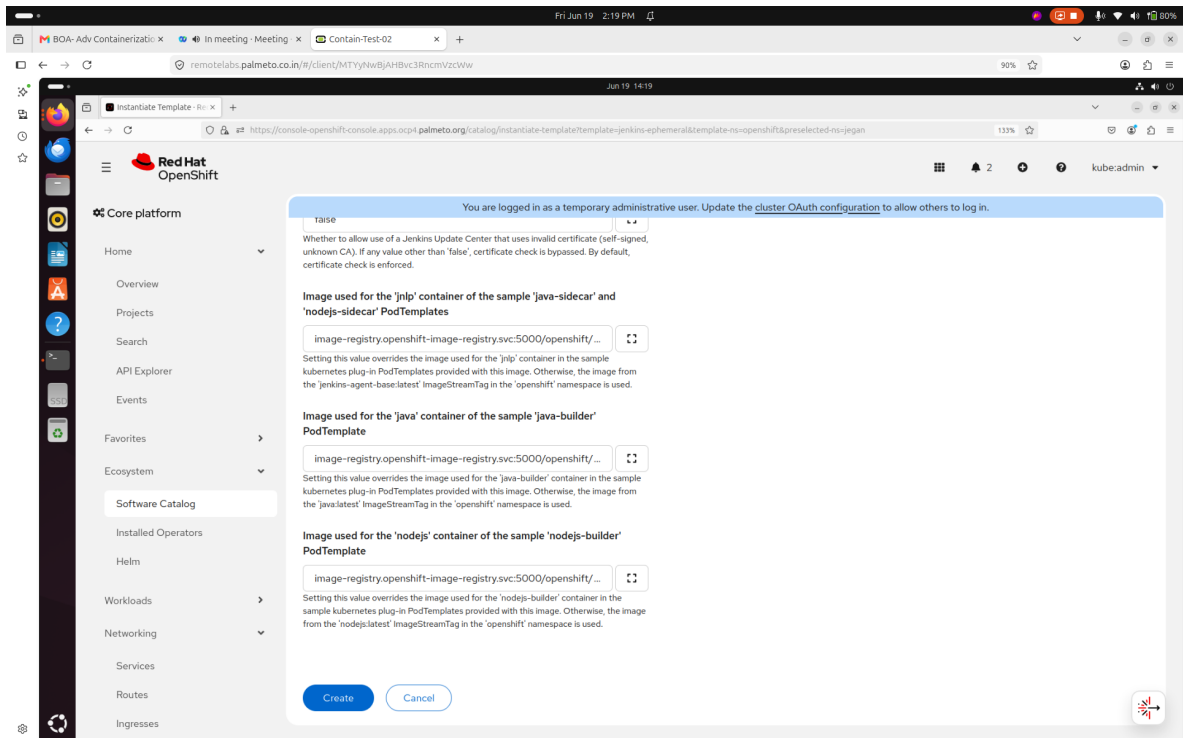
1 - 50 of 85

Name	Display name	Status	Requester	Memory	CPU	Created
abishek	No display name	Active	system:admin	471.1 MiB	0.001 cores	Jun 16, 2026, 4:19 PM
aishu	No display name	Active	kube:admin	915.2 MiB	0.004 cores	Jun 17, 2026, 3:03 PM
default	No display name	Active	No requester	66.9 MiB	-	Jun 14, 2026, 8:14 PM
dheena	No display name	Active	kube:admin	702.0 MiB	5.482 cores	5 minutes ago
franklin	No display name	Active	system:admin	360.1 MiB	0.001 cores	Jun 16, 2026, 3:52 PM
jms-demo	No display name	Active	system:admin	-	-	Jun 18, 2026, 5:23 PM
jms-dheena	No display name	Active	kube:admin	-	-	Jun 18, 2026, 5:23 PM
kube-node-lease	No display name	Active	No requester	-	-	Jun 14, 2026, 8:14 PM
kube-public	No display name	Active	No requester	-	-	Jun 14, 2026, 8:14 PM







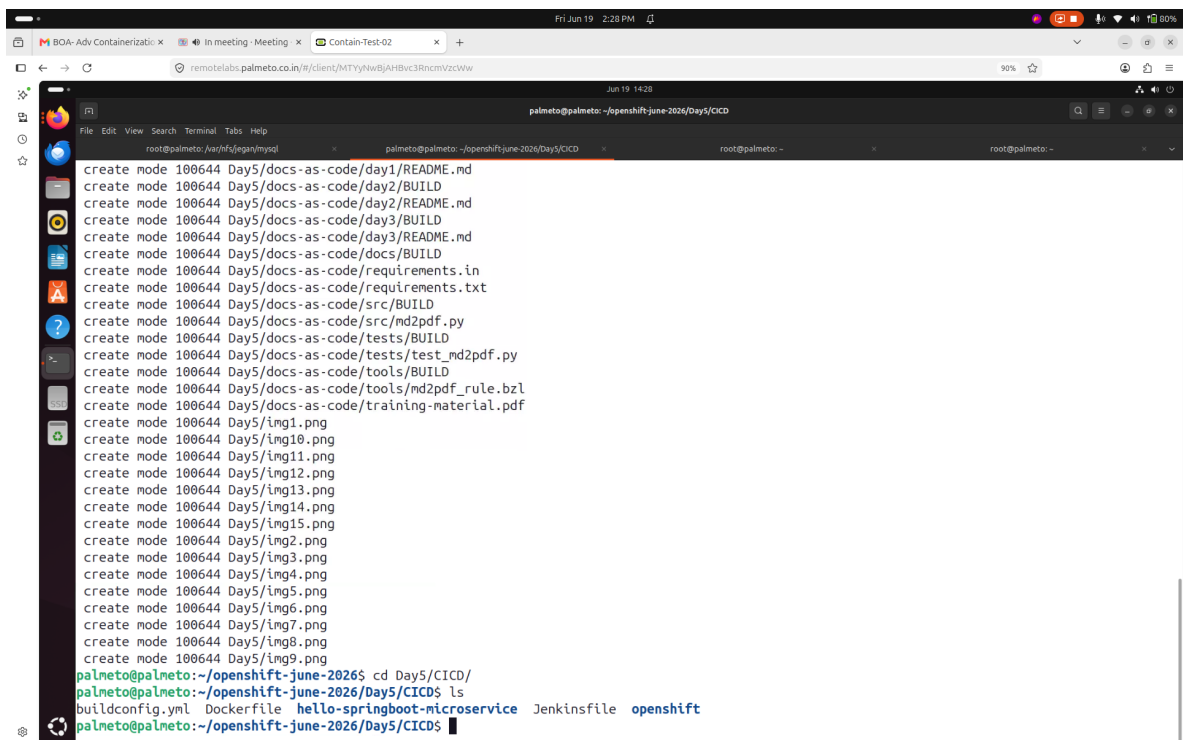





This screenshot shows the Red Hat OpenShift console interface. The left sidebar contains navigation options like Home, Overview, Projects, Search, API Explorer, Events, Favorites, Ecosystem, Software Catalog, Installed Operators, Helm, Workloads, Networking, Services, Routes, and Ingresses. The main panel displays the details for the 'jenkins' application in the 'jegan' project. A notification at the top states: 'You are logged in as a temporary administrative user. Update the cluster OAuth configuration to allow others to log in.' Below this, there's a search bar and filters. The application icon for Jenkins is shown. A warning box indicates: 'DeploymentConfig is being deprecated with OpenShift 4.14. Feature development of DeploymentConfigs will be deprecated in OpenShift Container Platform 4.14. DeploymentConfigs will continue to be supported for security and critical fixes, but you should migrate to Deployments wherever it is possible.' The 'Pods' section shows 'jenkins-1-sgqd8' in a 'Running' state. The 'Services' section shows 'jenkins-info' with a service port of 5000.

This screenshot shows the Red Hat OpenShift console interface, specifically the logs for the 'jenkins-1-sgqd8' pod. The left sidebar is the same as the previous screenshot. The main panel shows the 'jenkins-1-sgqd8' pod in a 'Running' state. The 'Logs' tab is selected, displaying a list of log entries. The logs show the Jenkins startup process, including the URL set to 'https://jenkins-jegan.apps.ocp4.palmeto.org/', the Java command used to start the application, and the successful initialization of the Jenkins server. The logs end with the message: '2026-06-19 08:50:42 INFO jenkins.InitReactorRunner\$1 onAttained Listed all plugins'.





```
BOA-Adv Containerizati... In meeting - Meeting x Contain-Test-02 x +
remotelabs.palmeto.co.in/ri/client/MTYyNWwBjAHBvc3RncmVzcWw
90%
Jun 19 14:28
palmeto@palmeto: ~/openshift-june-2026/Day5/CICD
File Edit View Search Terminal Tabs Help
root@palmeto: /ar/hfs/jegan/mysal palmeto@palmeto: ~/openshift-june-2026/Day5/CICD root@palmeto: - root@palmeto: -
create mode 100644 Day5/img14.png
create mode 100644 Day5/img15.png
create mode 100644 Day5/img2.png
create mode 100644 Day5/img3.png
create mode 100644 Day5/img4.png
create mode 100644 Day5/img5.png
create mode 100644 Day5/img6.png
create mode 100644 Day5/img7.png
create mode 100644 Day5/img8.png
create mode 100644 Day5/img9.png
palmeto@palmeto:~/openshift-june-2026$ cd Day5/CICD/
palmeto@palmeto:~/openshift-june-2026/Day5/CICD$ ls
buildconfig.yml Dockerfile hello-springboot-microservice Jenkinsfile openshift
palmeto@palmeto:~/openshift-june-2026/Day5/CICD$ oc project jegan
Already on project "jegan" on server "https://api.ocp4.palmeto.org:6443".
palmeto@palmeto:~/openshift-june-2026/Day5/CICD$ oc get all
Warning: apps.openshift.io/v1 DeploymentConfig is deprecated in v4.14+, unavailable no sooner than v6.0+
NAME READY STATUS RESTARTS AGE
pod/jenkins-1-deploy 0/1 Completed 0 8m25s
pod/jenkins-1-sgqd8 1/1 Running 0 8m24s

NAME DESIRED CURRENT READY AGE
replicationcontroller/jenkins-1 1 1 1 8m26s

NAME TYPE CLUSTER-IP EXTERNAL-IP PORT(S) AGE
service/jenkins ClusterIP 172.30.241.151 <none> 80/TCP 8m26s
service/jenkins-jnlp ClusterIP 172.30.30.77 <none> 50000/TCP 8m26s

NAME REVISION DESIRED CURRENT TRIGGERED BY
deploymentconfig.apps.openshift.io/jenkins 1 1 1 config,image(jenkins:2)

NAME HOST/PORT PATH SERVICES PORT TERMINATION WILDCARD
route.route.openshift.io/jenkins jenkins-jegan.apps.ocp4.palmeto.org jenkins <all> edge/Redirect None
palmeto@palmeto:~/openshift-june-2026/Day5/CICD$
```

```
BOA-Adv Containerizati... In meeting - Meeting x Contain-Test-02 x +
remotelabs.palmeto.co.in/ri/client/MTYyNWwBjAHBvc3RncmVzcWw
90%
Jun 19 14:29
palmeto@palmeto: ~/openshift-june-2026/Day5/CICD
File Edit View Search Terminal Tabs Help
root@palmeto: /ar/hfs/jegan/mysal palmeto@palmeto: ~/openshift-june-2026/Day5/CICD root@palmeto: - root@palmeto: -
Warning: apps.openshift.io/v1 DeploymentConfig is deprecated in v4.14+, unavailable no sooner than v6.0+
NAME READY STATUS RESTARTS AGE
pod/jenkins-1-deploy 0/1 Completed 0 8m25s
pod/jenkins-1-sgqd8 1/1 Running 0 8m24s

NAME DESIRED CURRENT READY AGE
replicationcontroller/jenkins-1 1 1 1 8m26s

NAME TYPE CLUSTER-IP EXTERNAL-IP PORT(S) AGE
service/jenkins ClusterIP 172.30.241.151 <none> 80/TCP 8m26s
service/jenkins-jnlp ClusterIP 172.30.30.77 <none> 50000/TCP 8m26s

NAME REVISION DESIRED CURRENT TRIGGERED BY
deploymentconfig.apps.openshift.io/jenkins 1 1 1 config,image(jenkins:2)

NAME HOST/PORT PATH SERVICES PORT TERMINATION WILDCARD
route.route.openshift.io/jenkins jenkins-jegan.apps.ocp4.palmeto.org jenkins <all> edge/Redirect None
palmeto@palmeto:~/openshift-june-2026/Day5/CICD$
palmeto@palmeto:~/openshift-june-2026/Day5/CICD$ oc create imagestream hello-microservice
imagestream.image.openshift.io/hello-microservice created
palmeto@palmeto:~/openshift-june-2026/Day5/CICD$
palmeto@palmeto:~/openshift-june-2026/Day5/CICD$ oc get is
NAME IMAGE REPOSITORY TAGS UPDATED
hello-microservice image-registry.openshift-image-registry.svc:5000/jegan/hello-microservice
palmeto@palmeto:~/openshift-june-2026/Day5/CICD$ oc describe is/hello-microservice
Name: hello-microservice
Namespace: jegan
Created: 21 seconds ago
Labels: <none>
Annotations: <none>
Image Repository: image-registry.openshift-image-registry.svc:5000/jegan/hello-microservice
Image Lookup: local=false
Tags: <none>
palmeto@palmeto:~/openshift-june-2026/Day5/CICD$
```



```

# This triggers the Jenkins Pipeline engine
strategy:
  type: JenkinsPipeline
  jenkinsPipelineStrategy:
    # Path to your Jenkinsfile in the Git repo
    jenkinsfilePath: Jenkinsfile

# Where your Jenkinsfile and Source Code live
source:
  type: Git
  git:
    uri: 'https://github.com/tektutor/openshift-june-2026.git'
    ref: 'main'
    contextDir: Day5/CICD

triggers:
- type: Generic
  generic:
    secretReference:
      name: generic-webhook-secret

palmeto@palmeto:~/openshift-june-2026/Day5/CICD$ oc policy add-role-to-user edit system:serviceaccount:jegan:default
clusterrole.rbac.authorization.k8s.io/edit added: "system:serviceaccount:jegan:default"
palmeto@palmeto:~/openshift-june-2026/Day5/CICD$ oc get buildconfig
No resources found in jegan namespace.
palmeto@palmeto:~/openshift-june-2026/Day5/CICD$ oc get builds
No resources found in jegan namespace.
palmeto@palmeto:~/openshift-june-2026/Day5/CICD$ oc apply -f buildconfig.yml
buildconfig.build.openshift.io/java-app-pipeline created
palmeto@palmeto:~/openshift-june-2026/Day5/CICD$ oc get buildconfig
NAME                                TYPE          FROM          LATEST
java-app-pipeline                   JenkinsPipeline  Git@main      0
palmeto@palmeto:~/openshift-june-2026/Day5/CICD$ oc get builds
No resources found in jegan namespace.
palmeto@palmeto:~/openshift-june-2026/Day5/CICD$
```

Jenkins Dashboard

Dashboard

+ New Item

People

Build History

Manage Jenkins

My Views

Open Blue Ocean

Lockable Resources

Build Queue

No builds in the queue.

Build Executor Status

1 idle

2 idle

3 idle

4 idle

5 idle

S	W	Name	Last Success	Last Failure	Last Duration	Fav
		jegan	N/A	N/A	N/A	

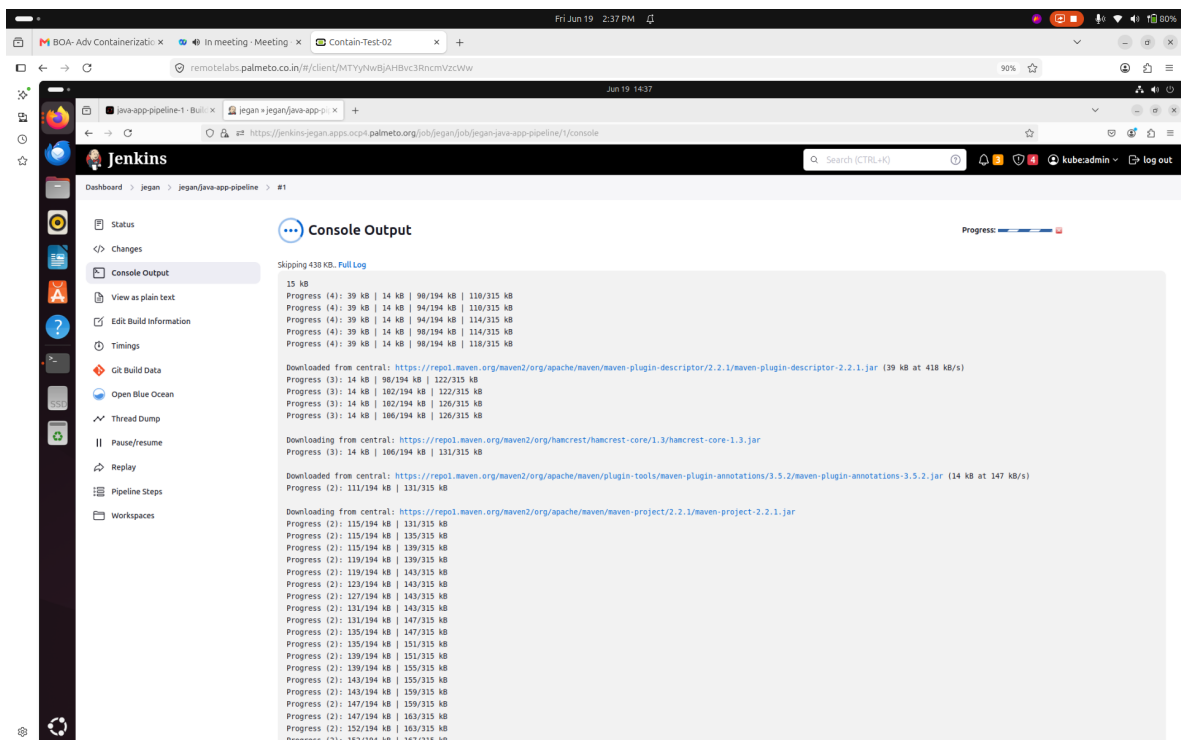

Icon legend



Atom feed for all

Atom feed for failures

Atom feed for just latest builds

REST API Jenkins 2.401.1





```
cd ~/openshift-july-2026
git pull
cd Day5/CICD
oc project jegan-project
oc apply -f buildconfig.yml
oc get buildconfigs
```

```
oc create imagestream hello-microservice
oc policy add-role-to-user edit system:serviceaccount:jegan-project:default
oc start-build bc/java-app-pipeline

oc logs -f bc/java-app-pipeline
```

Certifications Recommended

Recommended for Developers

- To prepare for EX288 Certification, you may attend the training that covers topics listed in
Containers & Kubernetes Fundamentals (D0188) and Red Hat OpenShift Development II: Building Kubernetes Applications (D0288).
- Red Hat doesn't mandate attending training before taking EX288 Certification, those who are technically hands-on will be able to clear the Certification with consistent preparation

- Red Hat Certified Specialist in Containers (EX188)
- Red Hat Certified Specialist in OpenShift Application Development (EX288)
- Red Hat Certified Cloud-Native Developer (EX378) - Quarkus/Java Focused
- Red Hat Certified Specialist in OpenShift AI (EX267)
- Event-Driven Development (EX453) - Kafka/AMQ Streams
- Building Resilient Microservices EX328 - Service Mesh/Istio

Recommended for DevOps Engineers

- Red Hat Certified Specialist in OpenShift Automation and Integration (EX380)
- Red Hat Certified Cloud-Native Developer (RHCCD)
- Red Hat Certified OpenShift Architect
- The Foundation: Red Hat Certified Engineer (EX294 - RHCE)
- The Platform: EX288 (Developer) OR EX280 (Admin)
- The Pipeline: EX288 (OpenShift Pipelines)
- The Automation Pinnacle: Red Hat Certified Specialist in MultiCluster Management (EX480)

Recommended for Administrators

- Prerequisite: Red Hat Certified System Administrator (EX200 - RHCSA)
- Core: Red Hat Certified Specialist in OpenShift Administration (EX280)
- Advanced: Red Hat Certified Specialist in OpenShift Automation & Integration (EX380)
- OpenShift Virtualization (EX316) - Running VMs alongside containers
- OpenShift Data Foundation (EX370) - Managing cluster storage/ODF